



# **Persistence II: File System Implementation, RAID, and InfiniCache**

*CS 571: Operating Systems (Spring 2022)*

Lecture 9

Yue Cheng

Some material taken/derived from:

- Wisconsin CS-537 materials by Remzi Arpaci-Dusseau.

Licensed for use under a Creative Commons Attribution-NonCommercial-ShareAlike 3.0 Unported License.

# **File System Implementation**

# File System Implementation

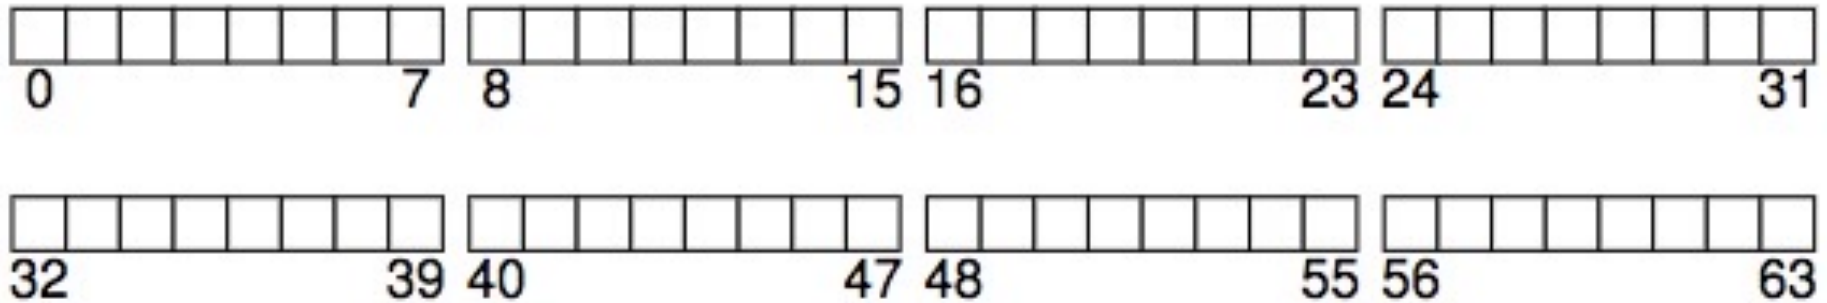
- On-disk structures
  - How do we represent files and directories?
- File system operations (internally)
  - How on-disk structures get touched when performing FS operations
- File system locality & data layout policies
  - How data layout impacts locality for on-disk FS?

# On-Disk Structures

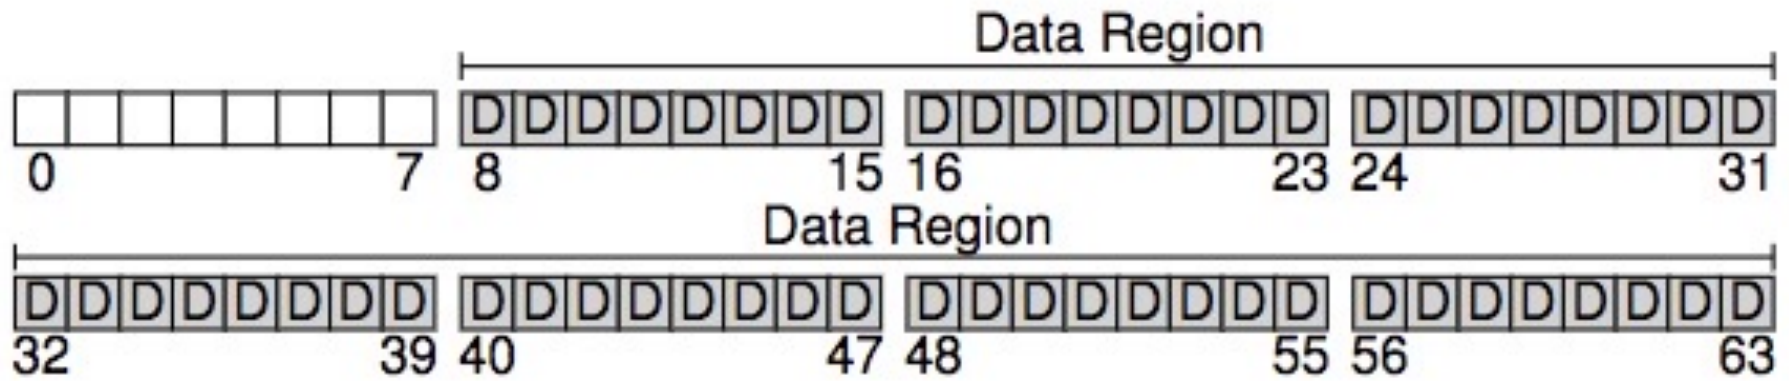
# On-Disk Structures

- Common file system structures
  - Data block
  - inode table
  - Directories
  - Data bitmap
  - inode bitmap
  - Superblock

# On-Disk Structure: Empty Disk



# On-Disk Structure: Data Blocks

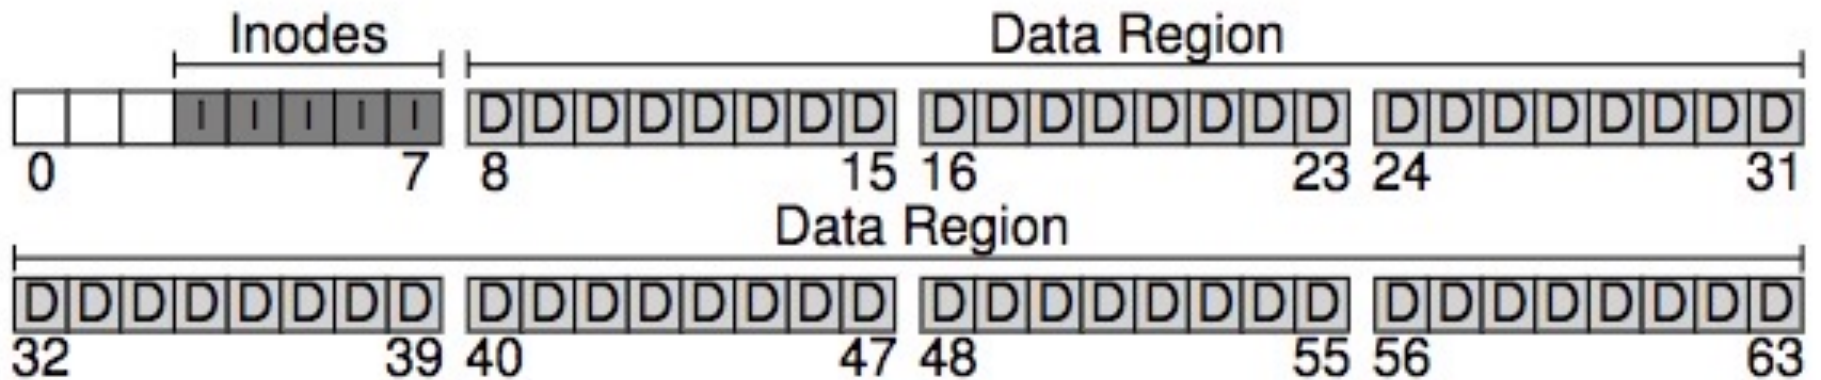


# On-Disk Structures

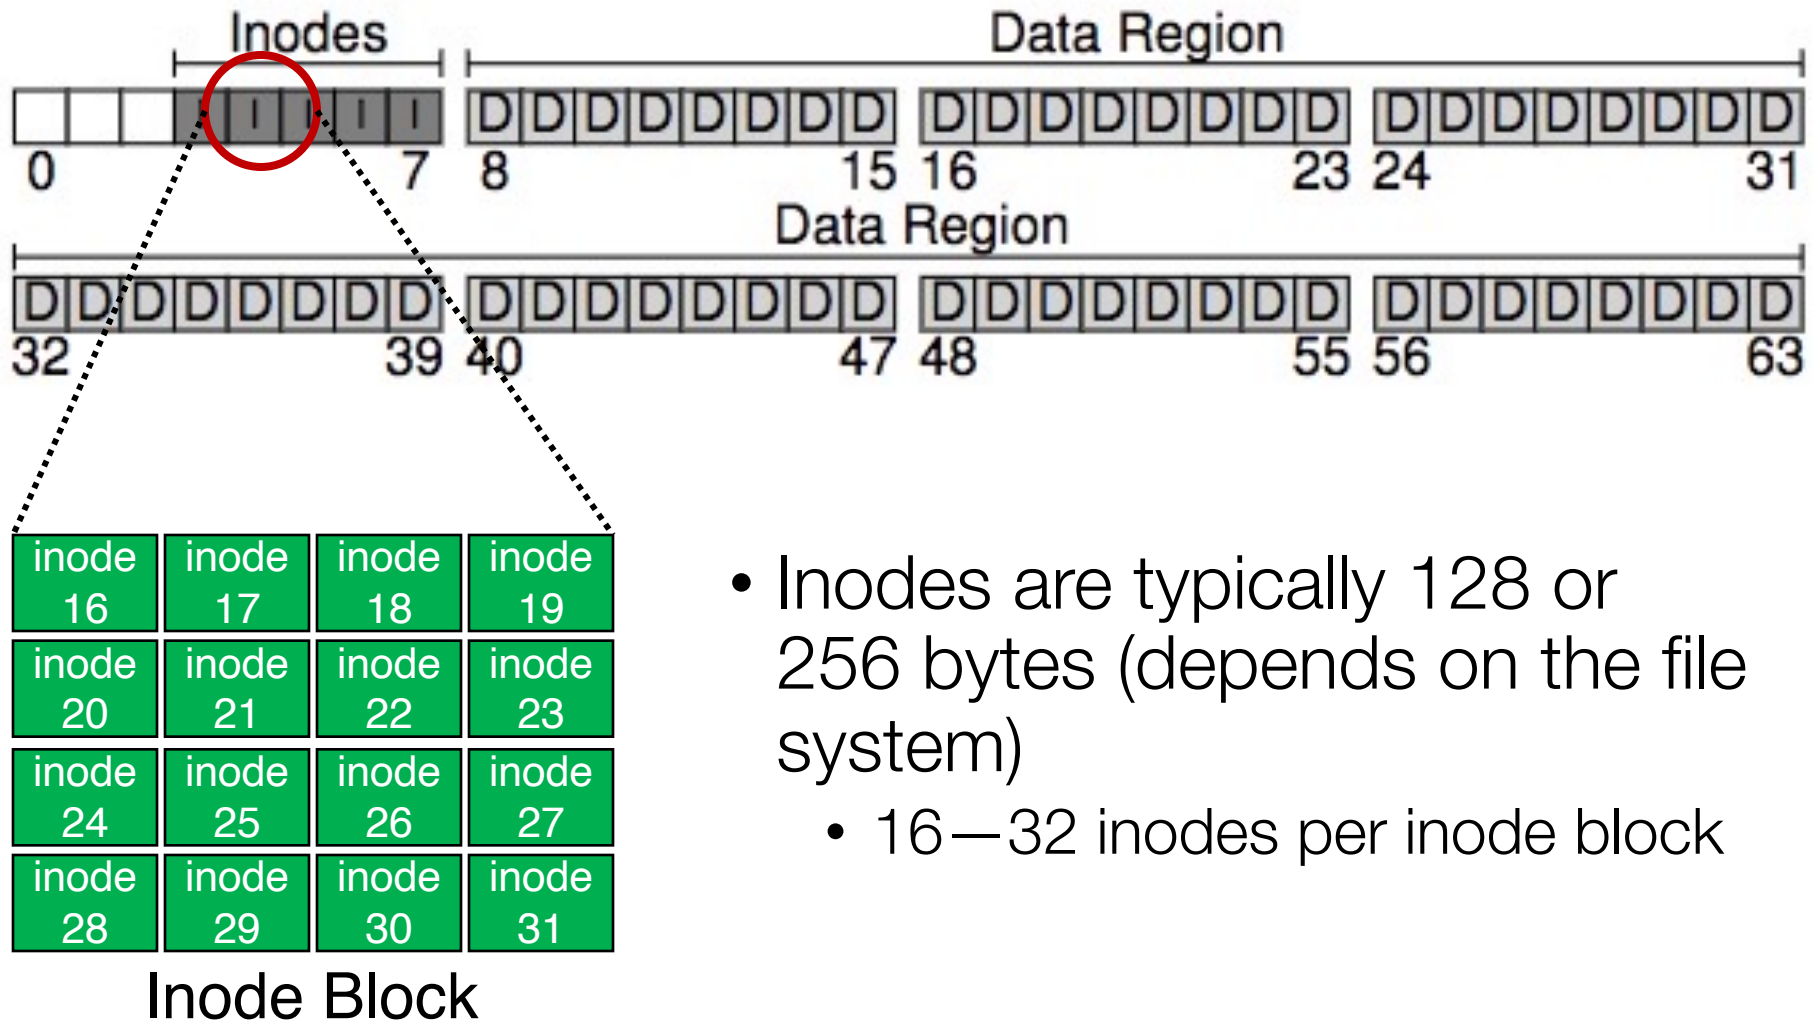
- Common file system structures
  - Data block
  - **inode table**
  - Directories
  - Data bitmap
  - inode bitmap
  - Superblock



# On-Disk Structure: Inodes

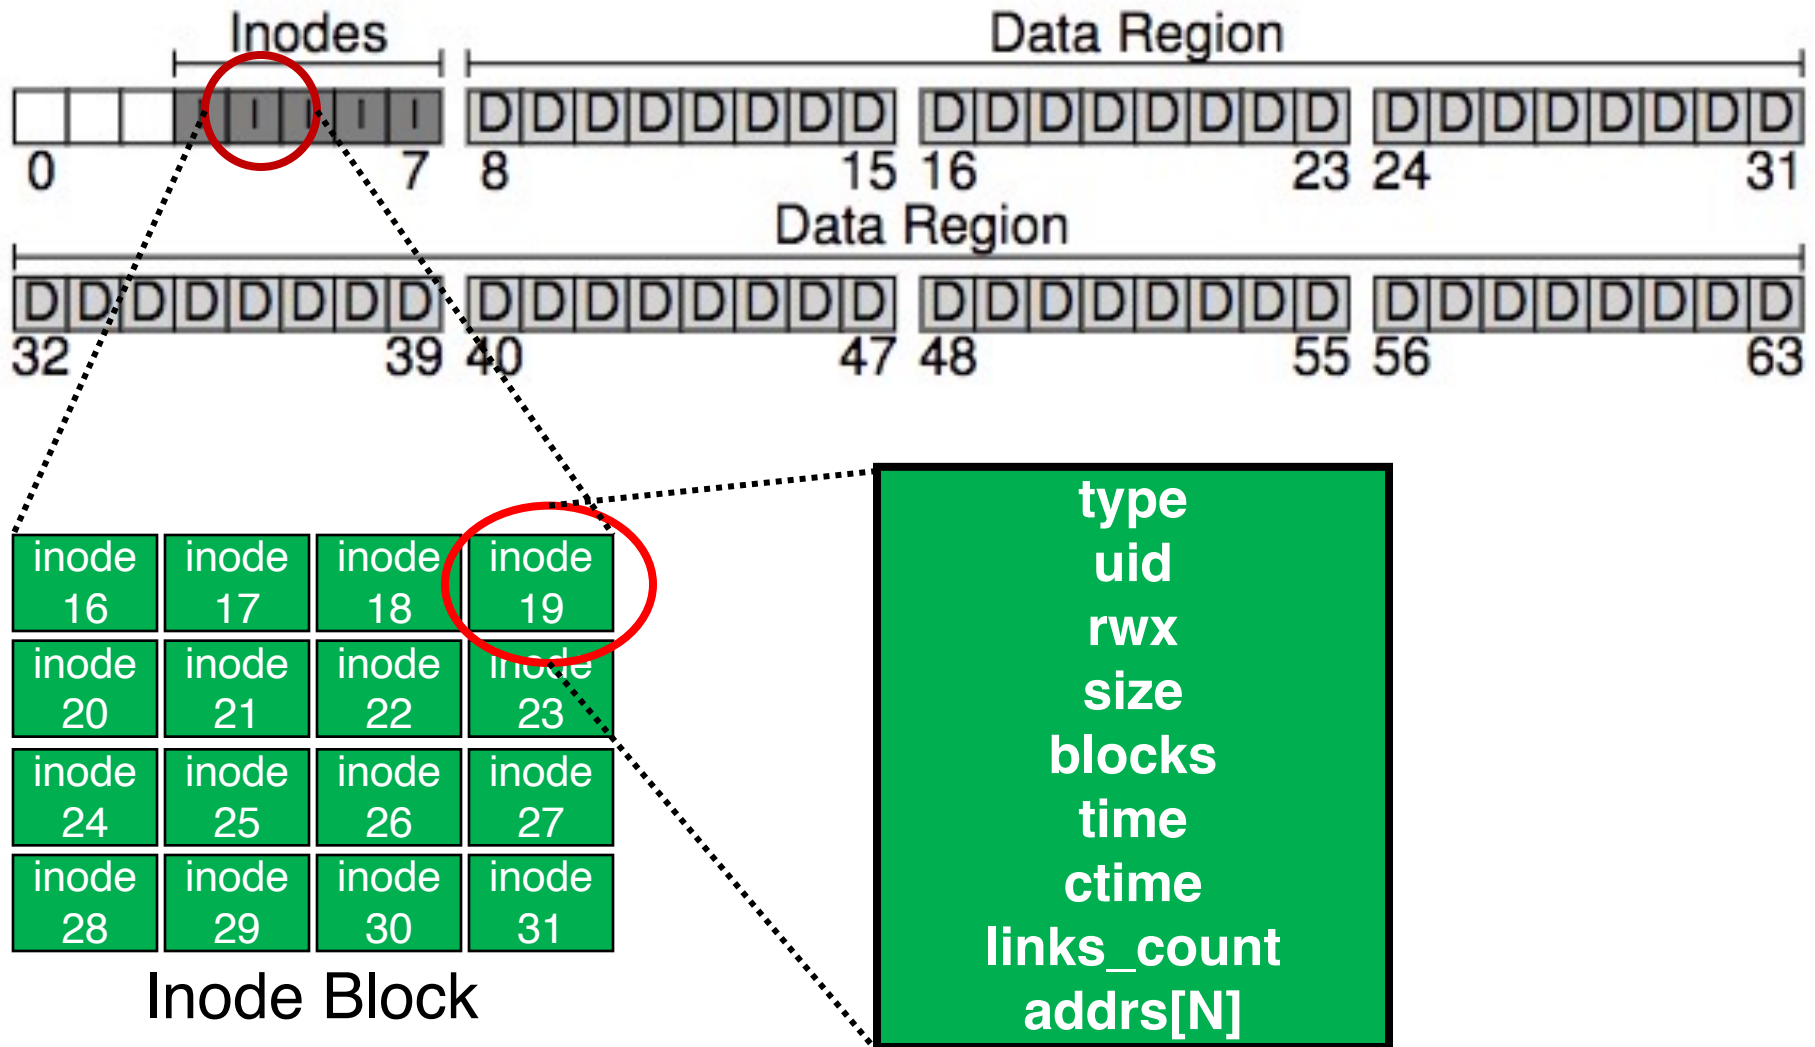


# On-Disk Structure: Inodes

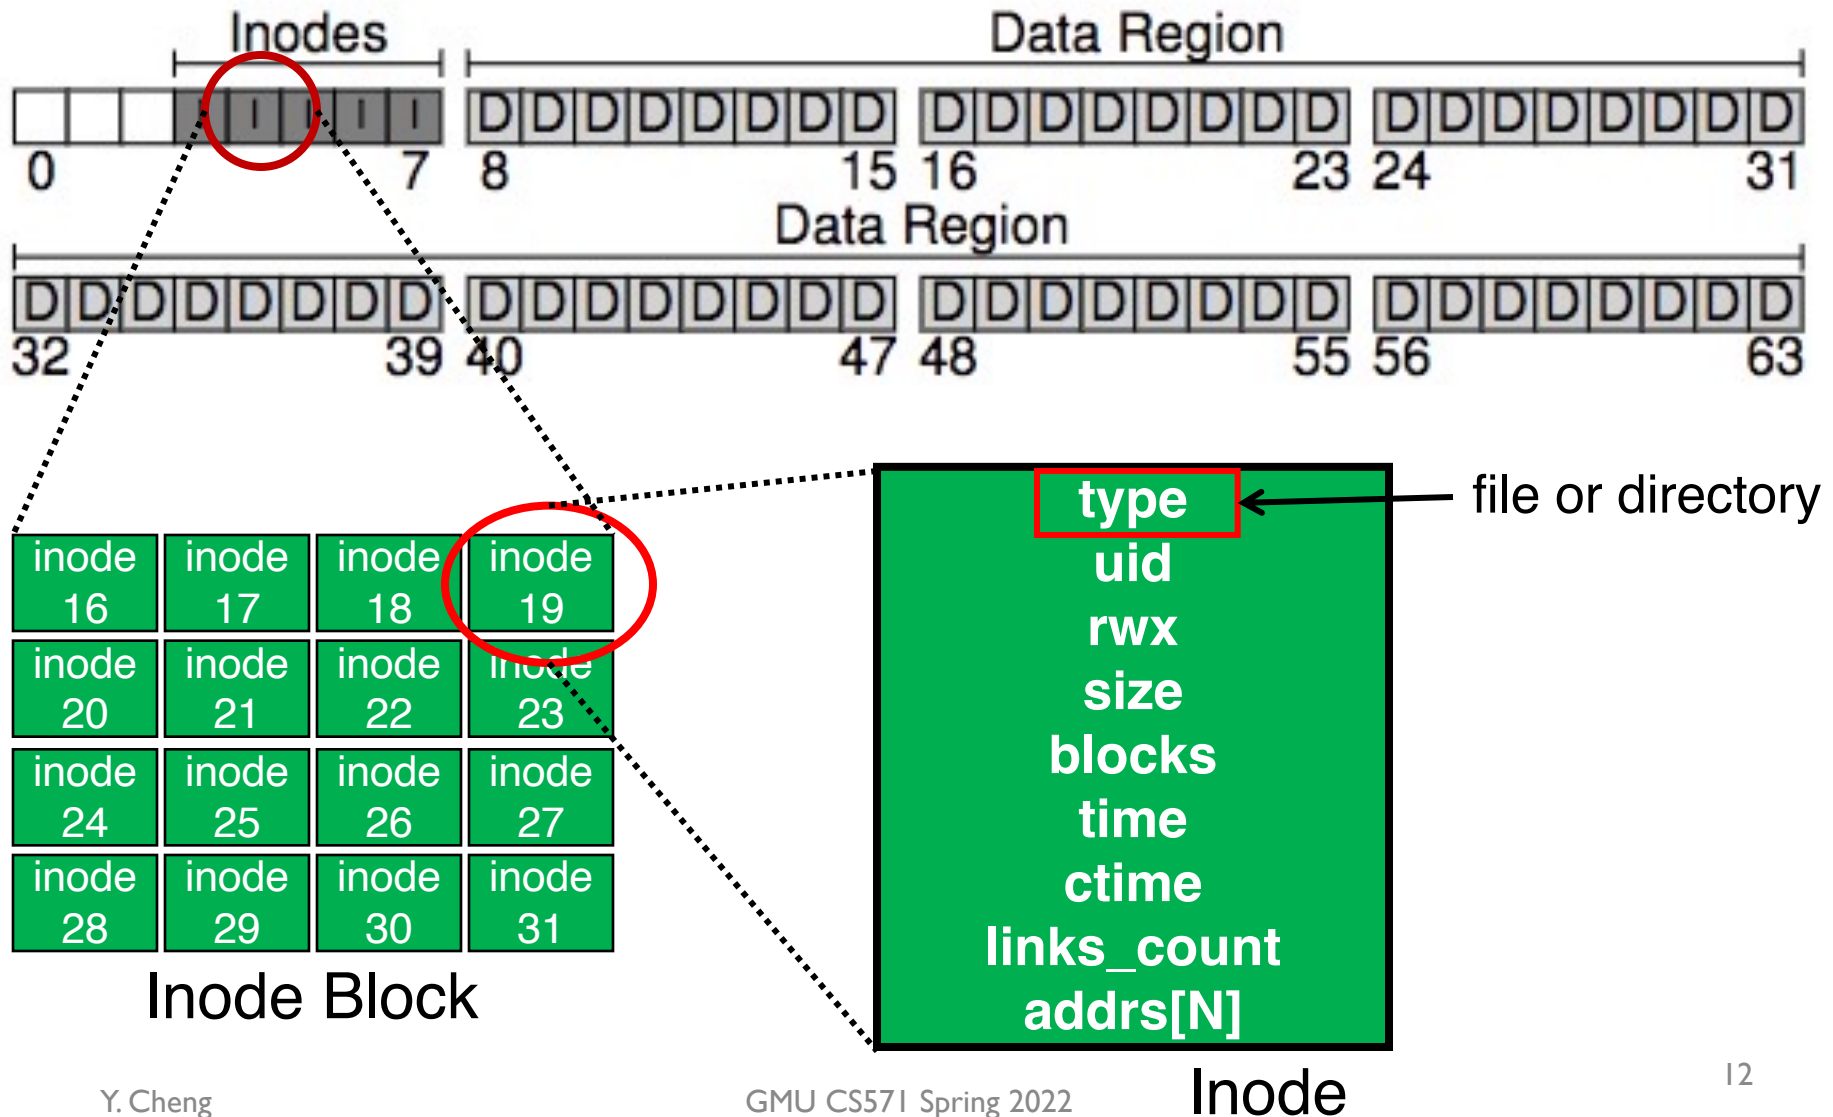


- Inodes are typically 128 or 256 bytes (depends on the file system)
  - 16—32 inodes per inode block

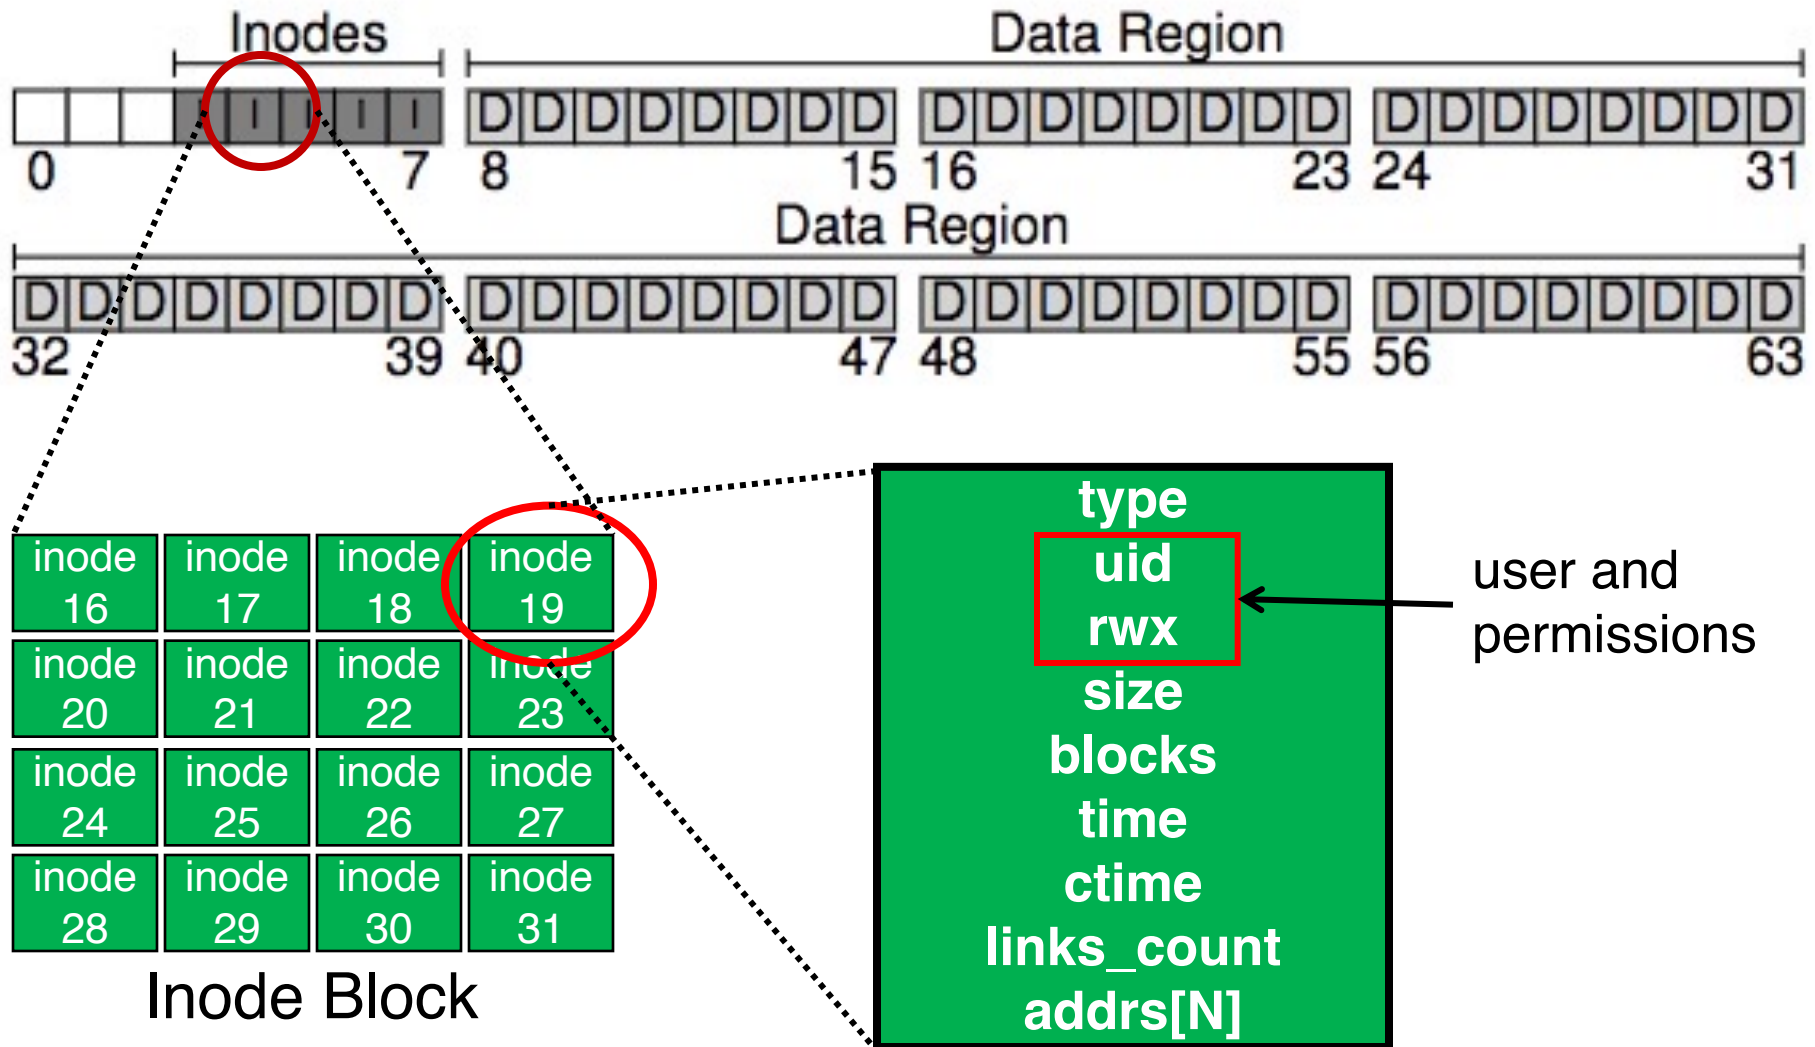
# On-Disk Structure: Inodes



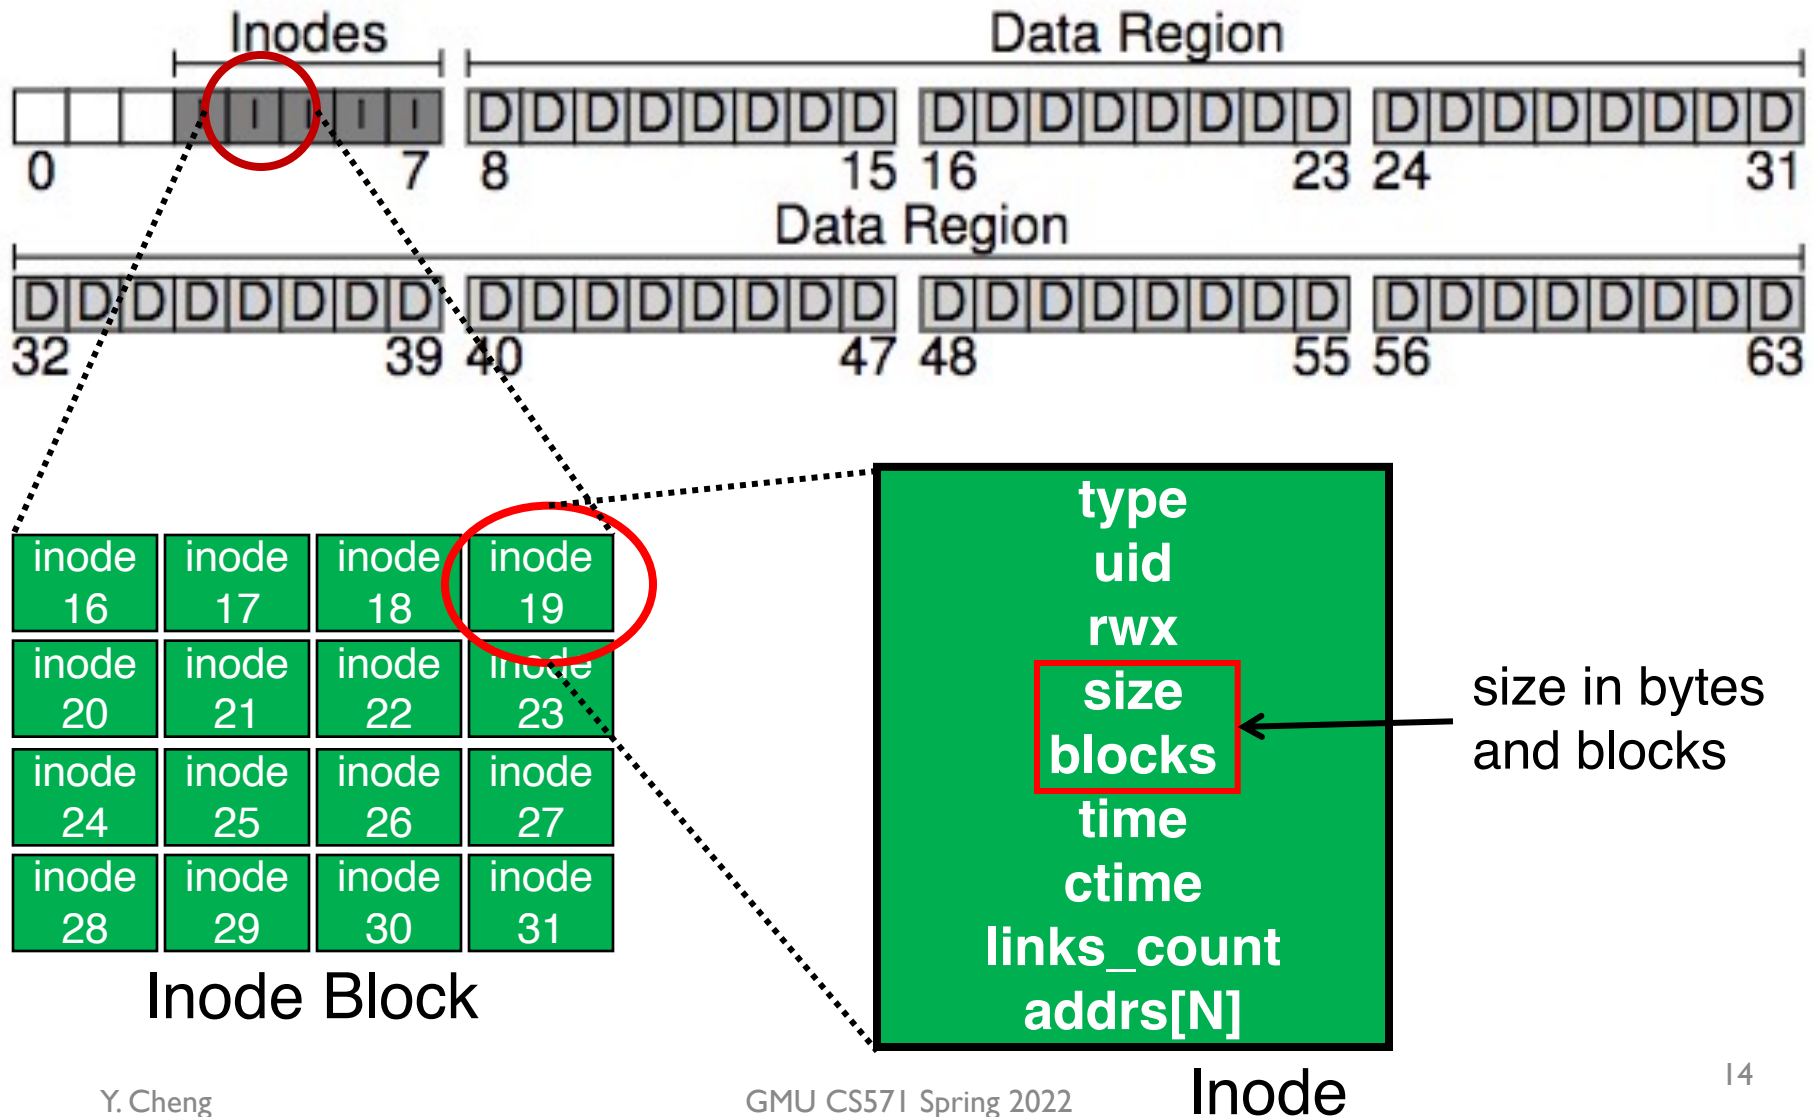
# On-Disk Structure: Inodes



# On-Disk Structure: Inodes

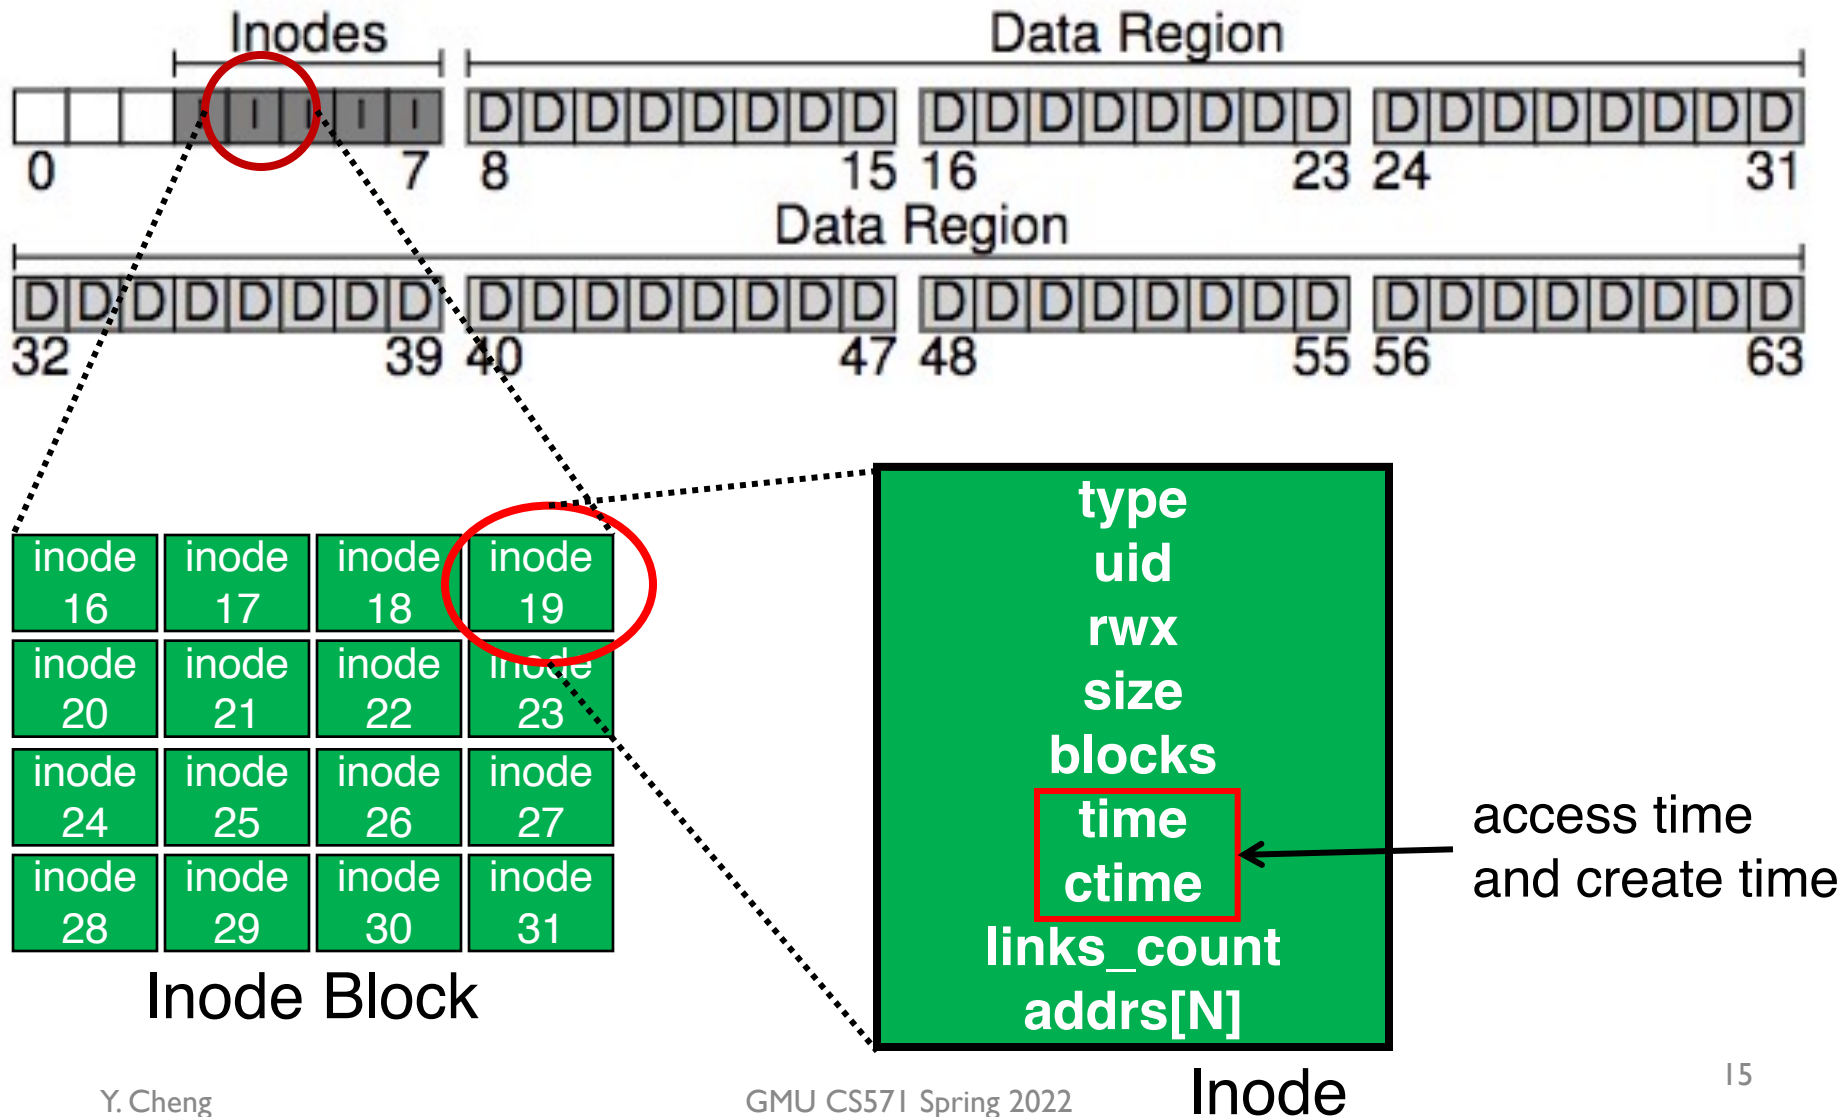


# On-Disk Structure: Inodes

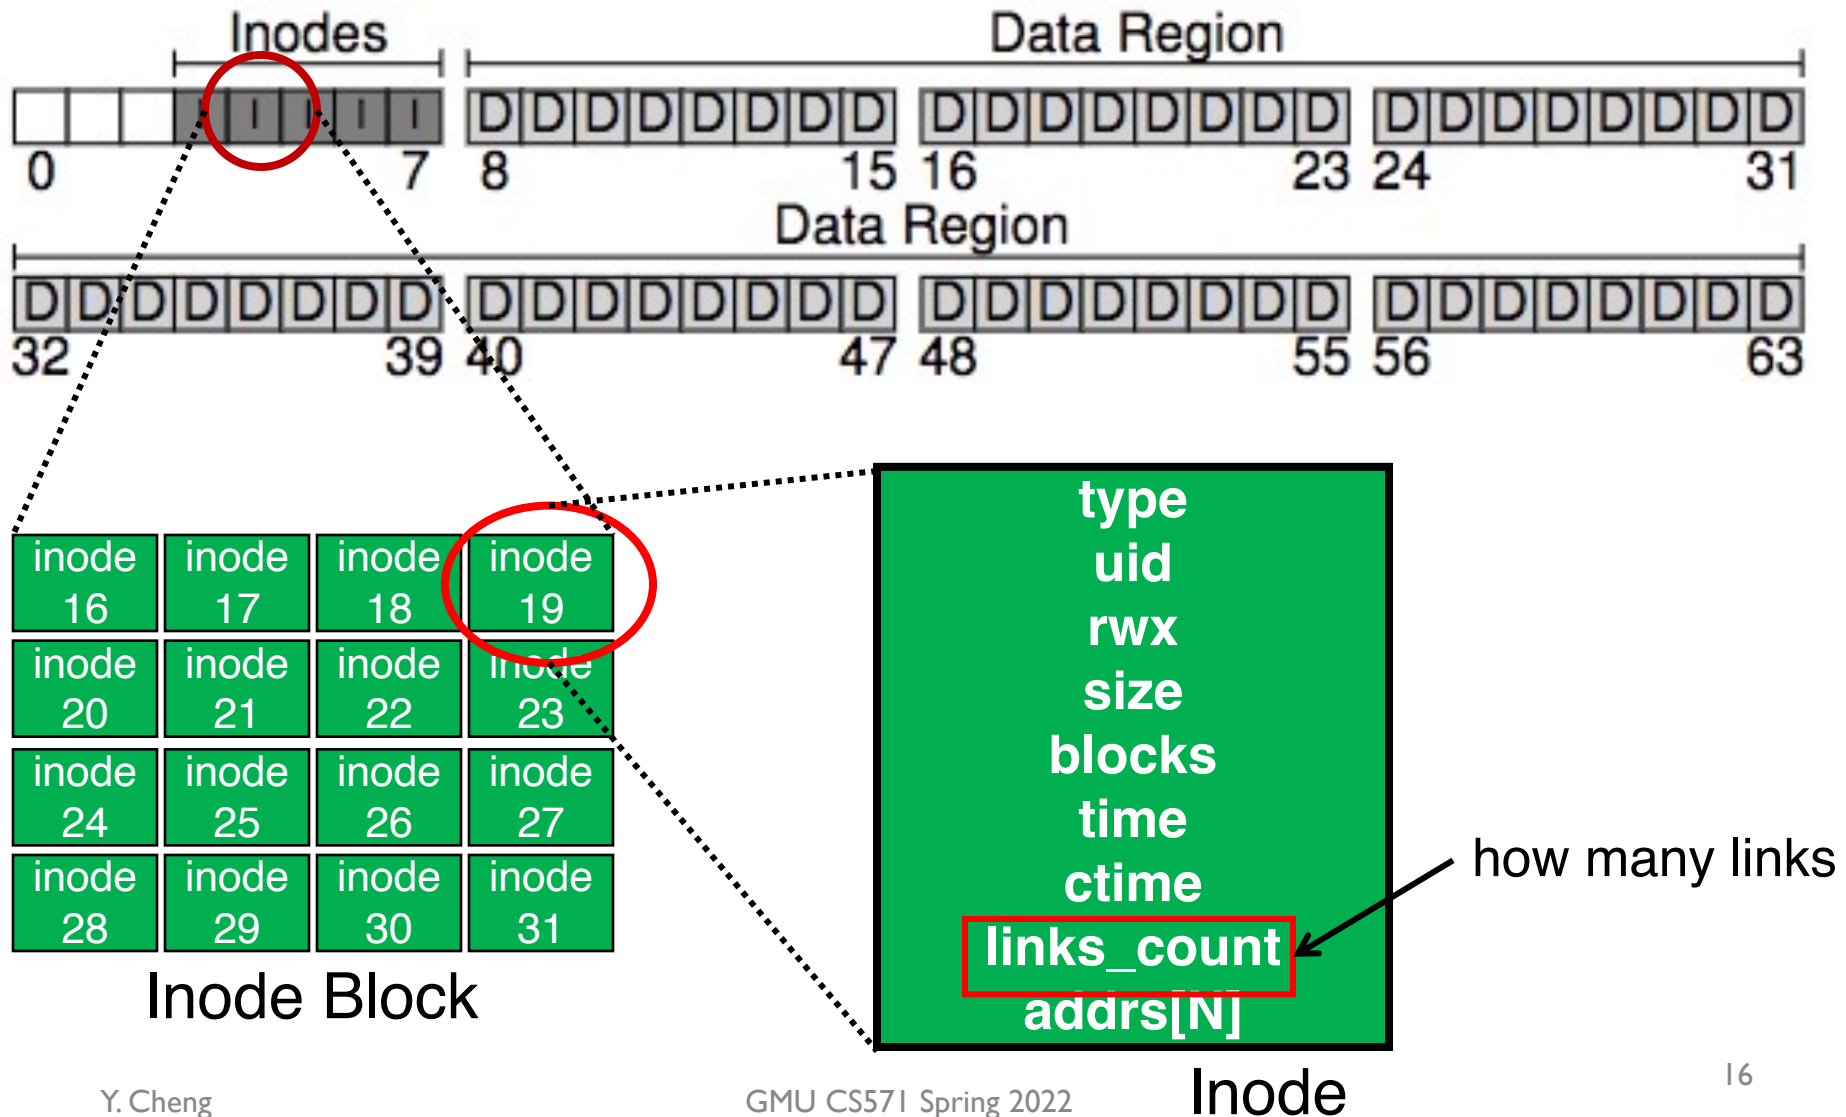




# On-Disk Structure: Inodes

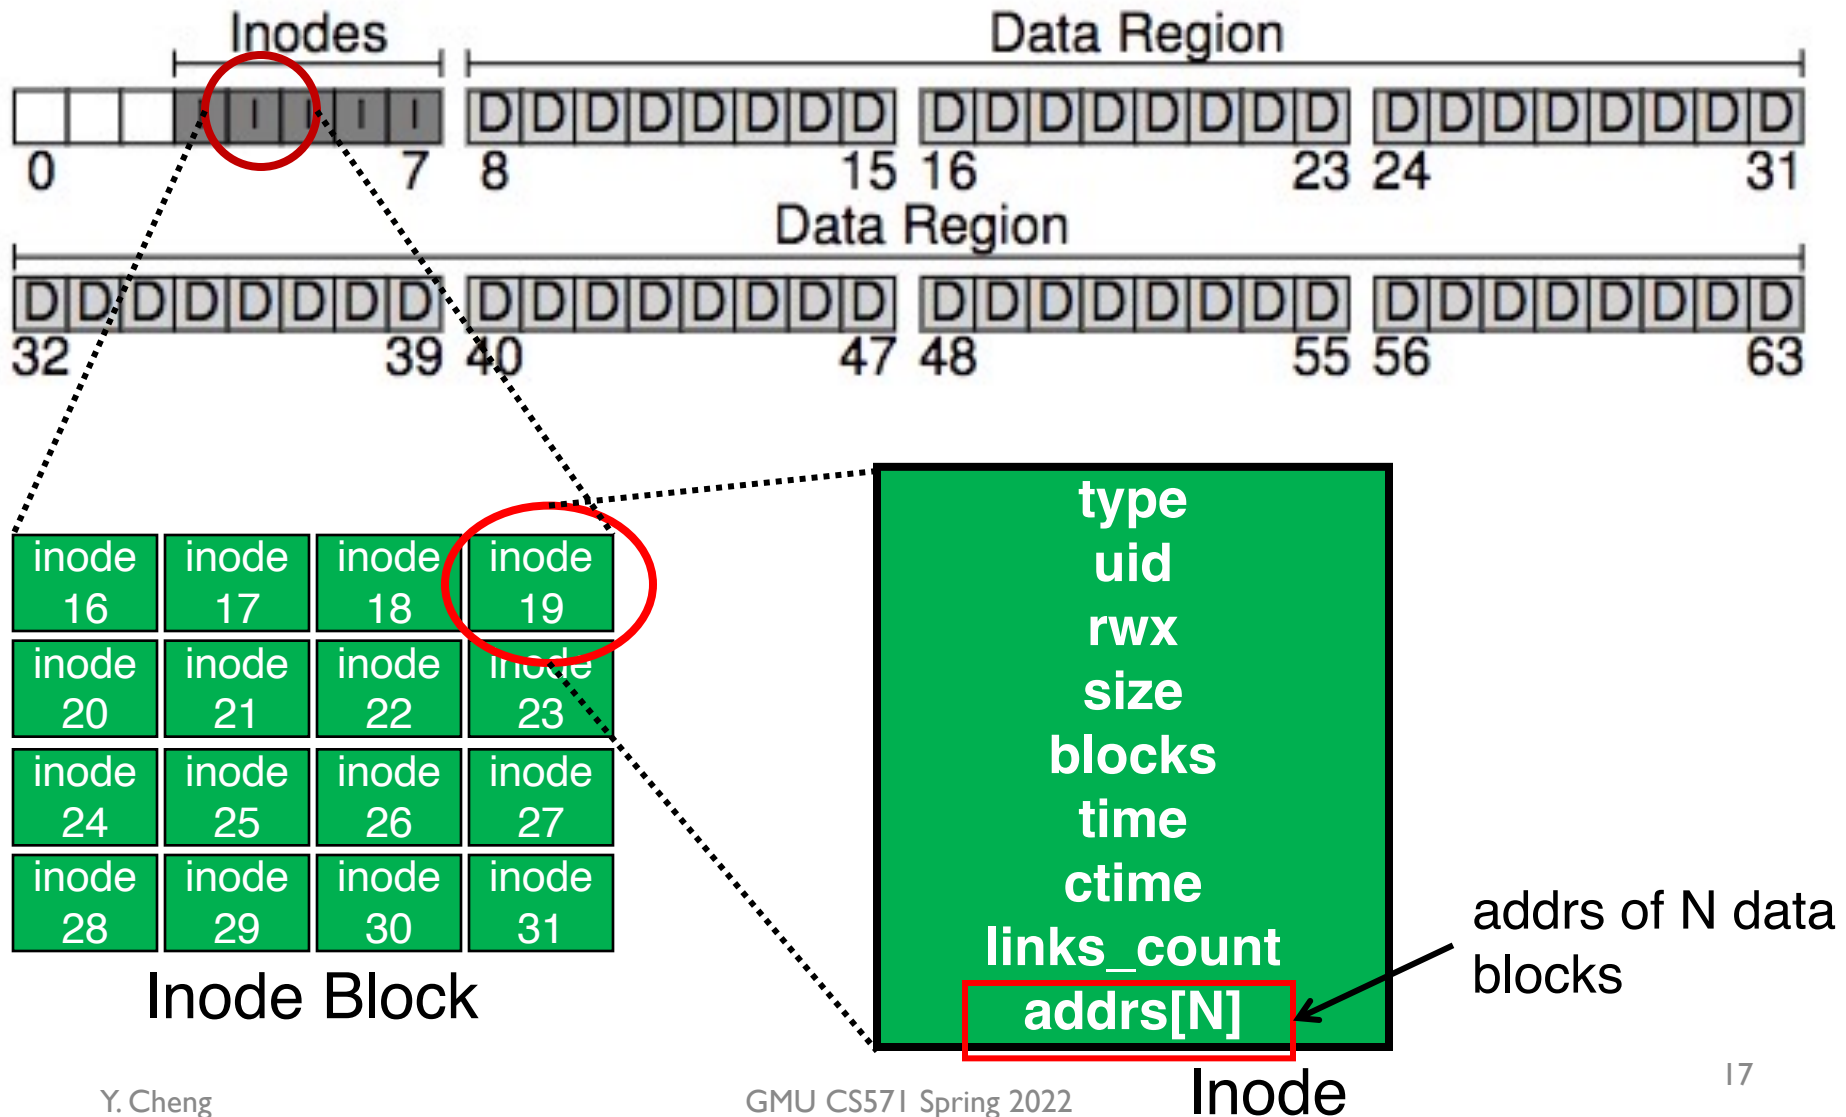


# On-Disk Structure: Inodes

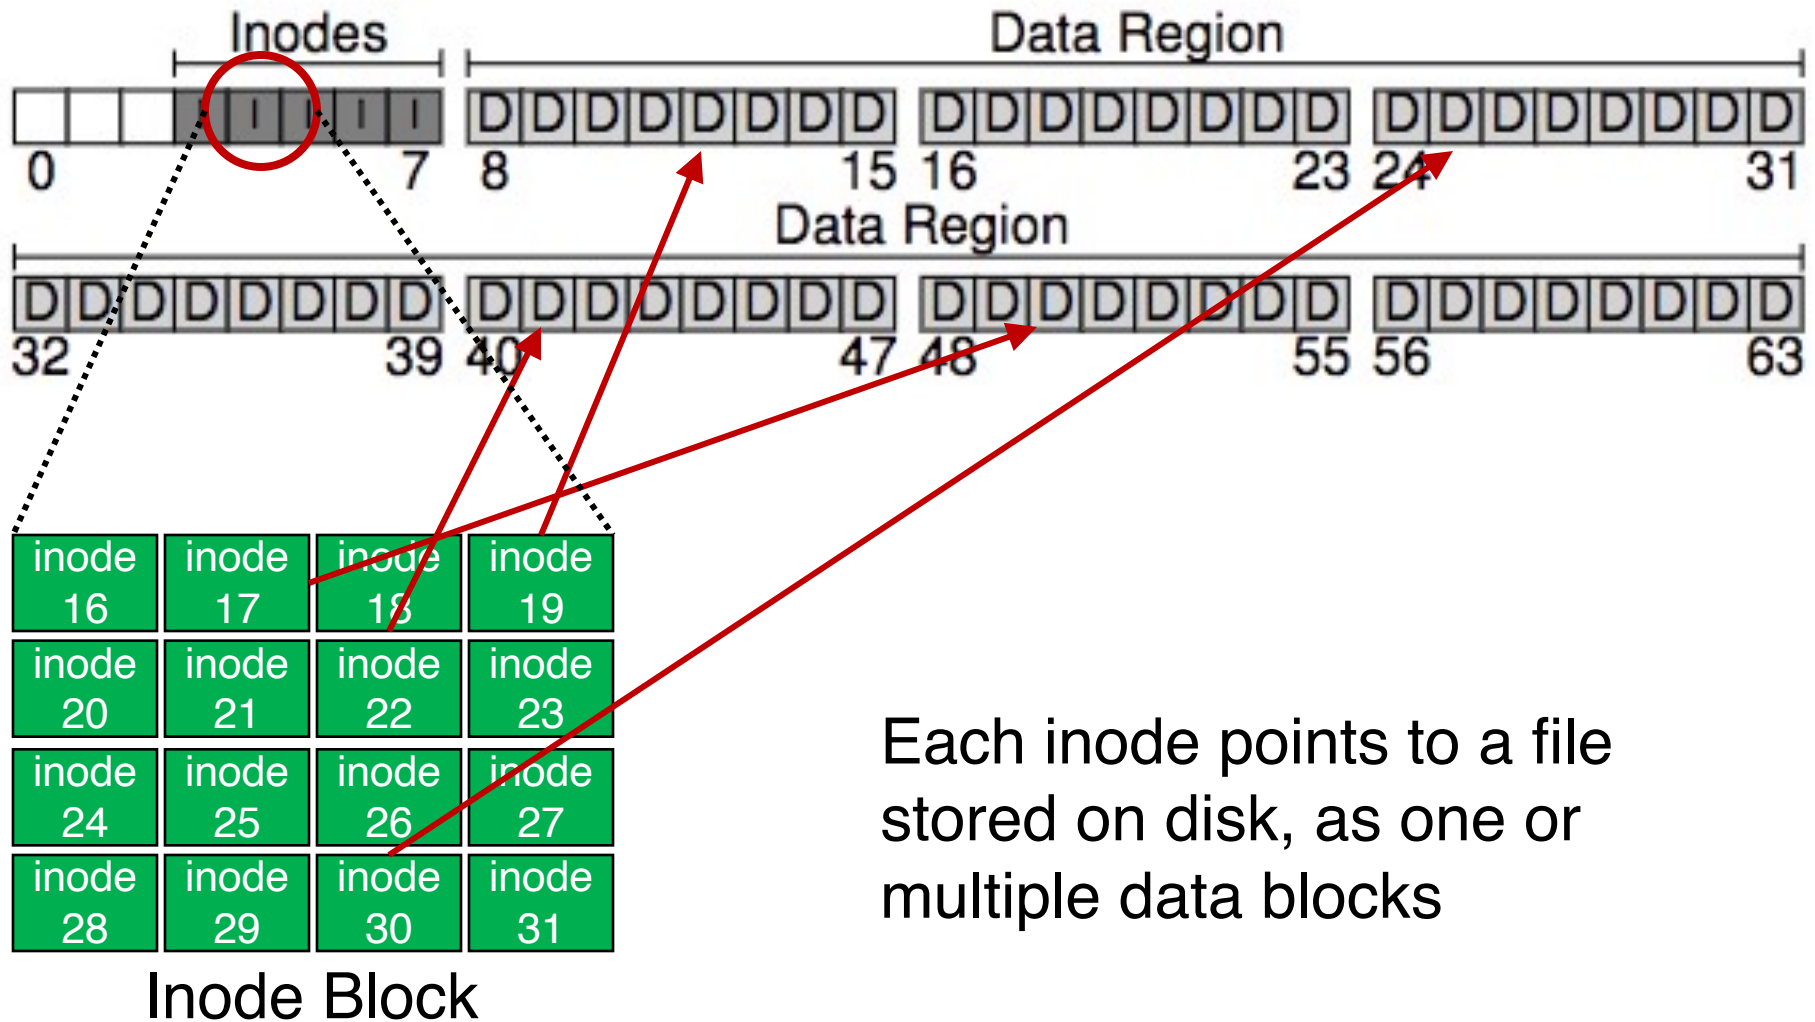




# On-Disk Structure: Inodes



# On-Disk Structure: Inodes



# On-Disk Structures

- Common file system structures
  - Data block
  - Inode table
  - Directories
  - Data bitmap
  - Inode bitmap
  - Superblock

# On-Disk Structure: Directories

- Common directory design: just store directory entries in files
  - Different file systems vary
- Various data structures (formats) could be used
  - Lists
  - B-trees

# On-Disk Structures

- Common file system structures
  - Data block
  - inode table
  - Directories
  - Data bitmap
  - inode bitmap
  - Superblock

# Allocation

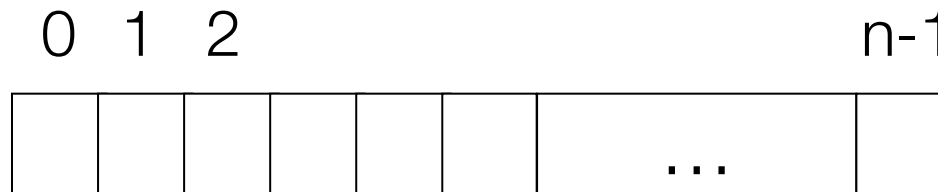
- How does file system find free data blocks or free inodes?

# Allocation

- How does file system find free data blocks or free inodes?
  - Free list
  - Bitmaps
- What are the tradeoffs?

# Bitmap

Each bit of the bitmap is used to indicate whether the corresponding object/block is **free** (0) or **in-use** (1)



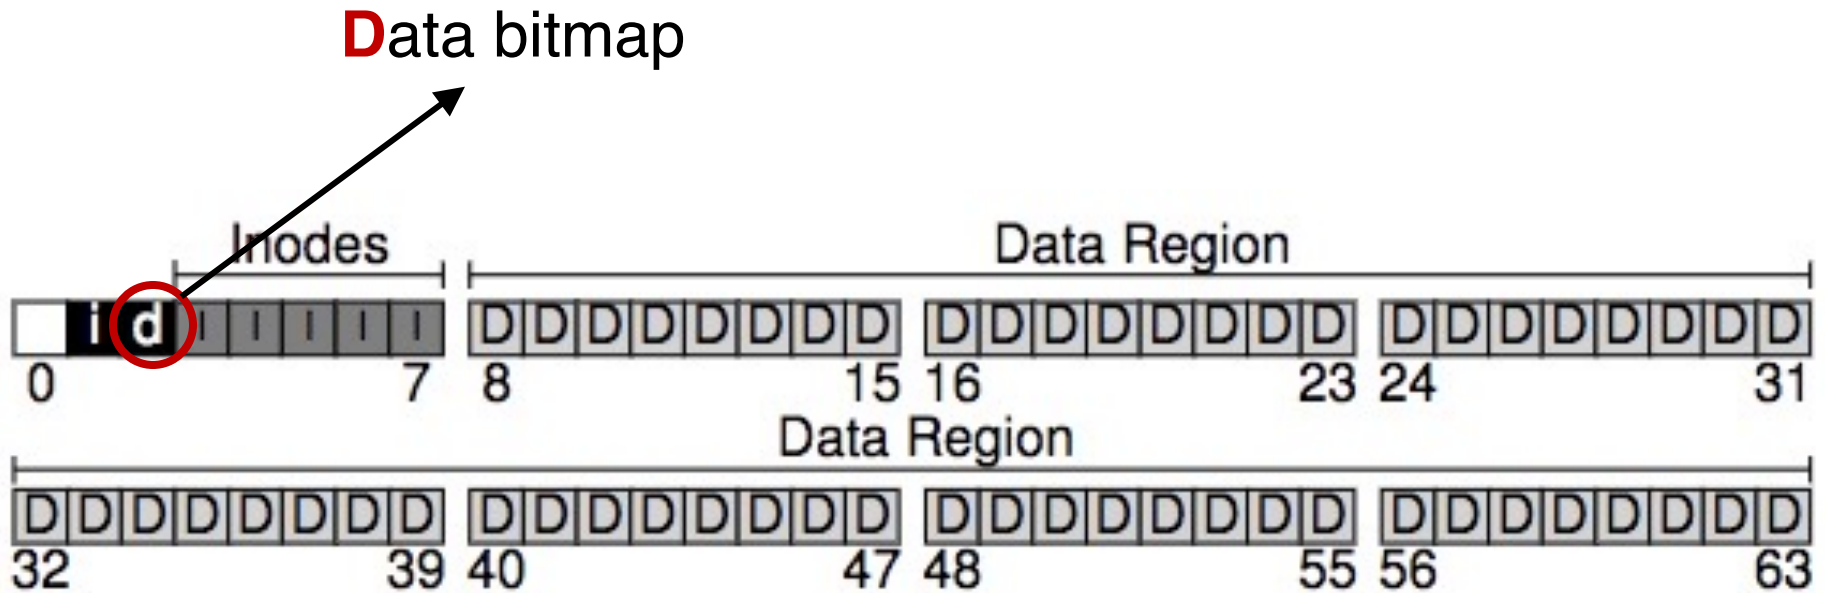
$$\text{bit}[i] = \begin{cases} 1 \Rightarrow \text{object}[i] \text{ in use} \\ 0 \Rightarrow \text{object}[i] \text{ free} \end{cases}$$



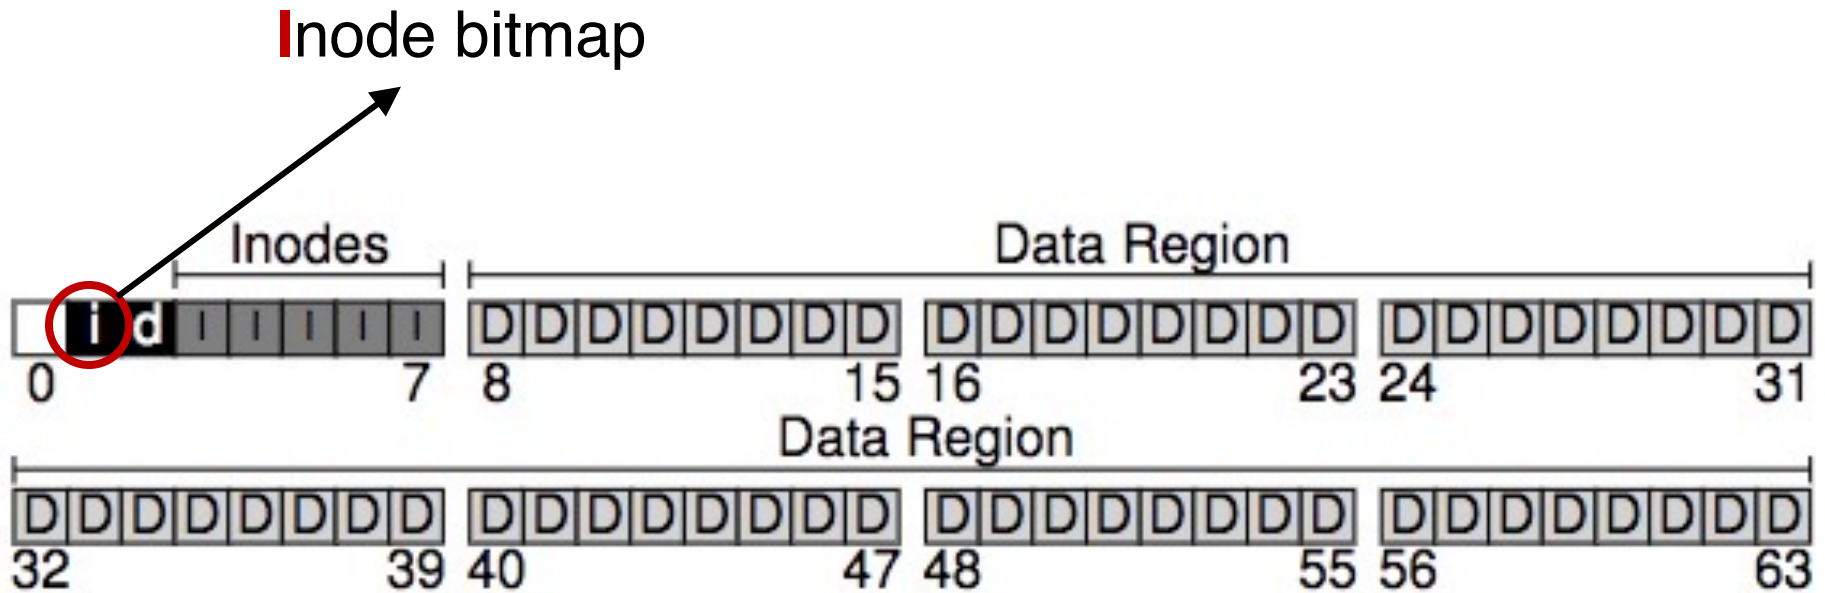
# Allocation

- How does file system find free data blocks or free inodes?
  - Free list
  - Bitmaps
- What are the tradeoffs?
  - Free list: Cannot get contiguous space easily
  - Bitmap: Easy to allocate contiguous space for files

# On-Disk Structure: Data Bitmaps



# On-Disk Structure: Inode Bitmaps



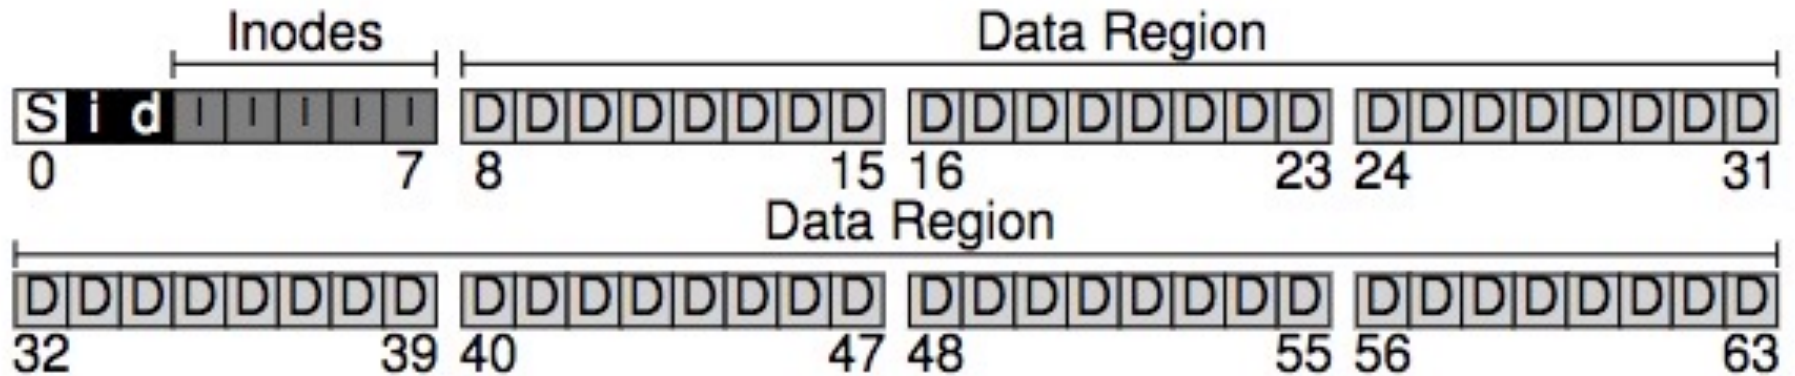
# On-Disk Structures

- Common file system structures
  - Data block
  - Inode table
  - Directories
  - Data bitmap
  - Inode bitmap
  - Superblock

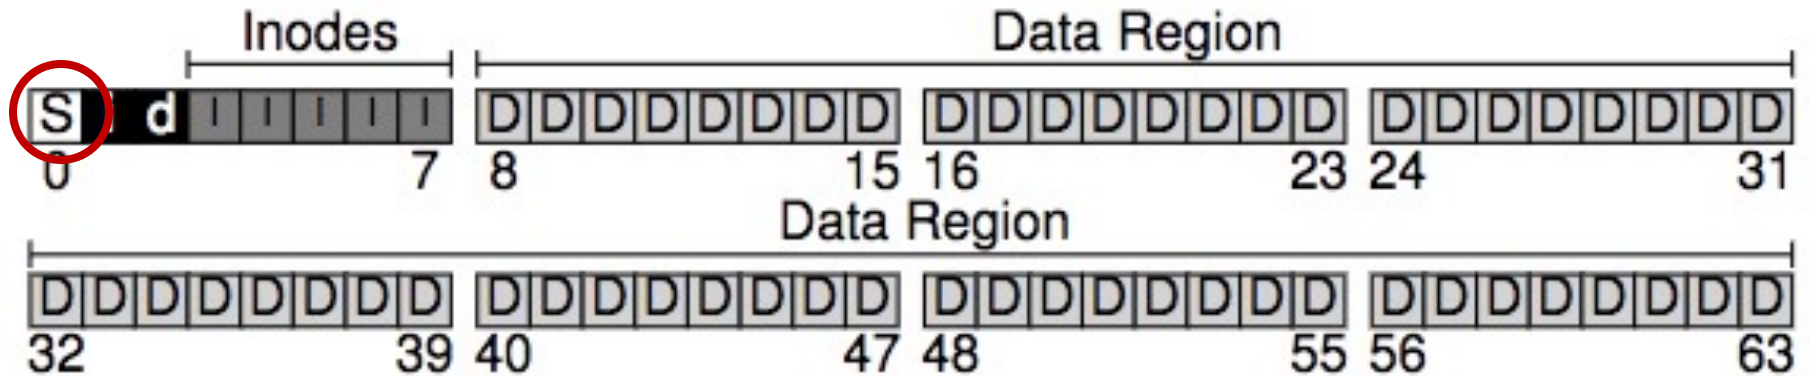
# On-Disk Structure: Superblock

- Need to know basic file system configuration and runtime status, such as:
  - Block size
  - How many inodes are there
  - How much free space
- Store all these **metadata** info in a superblock

# On-Disk Structure: Superblock

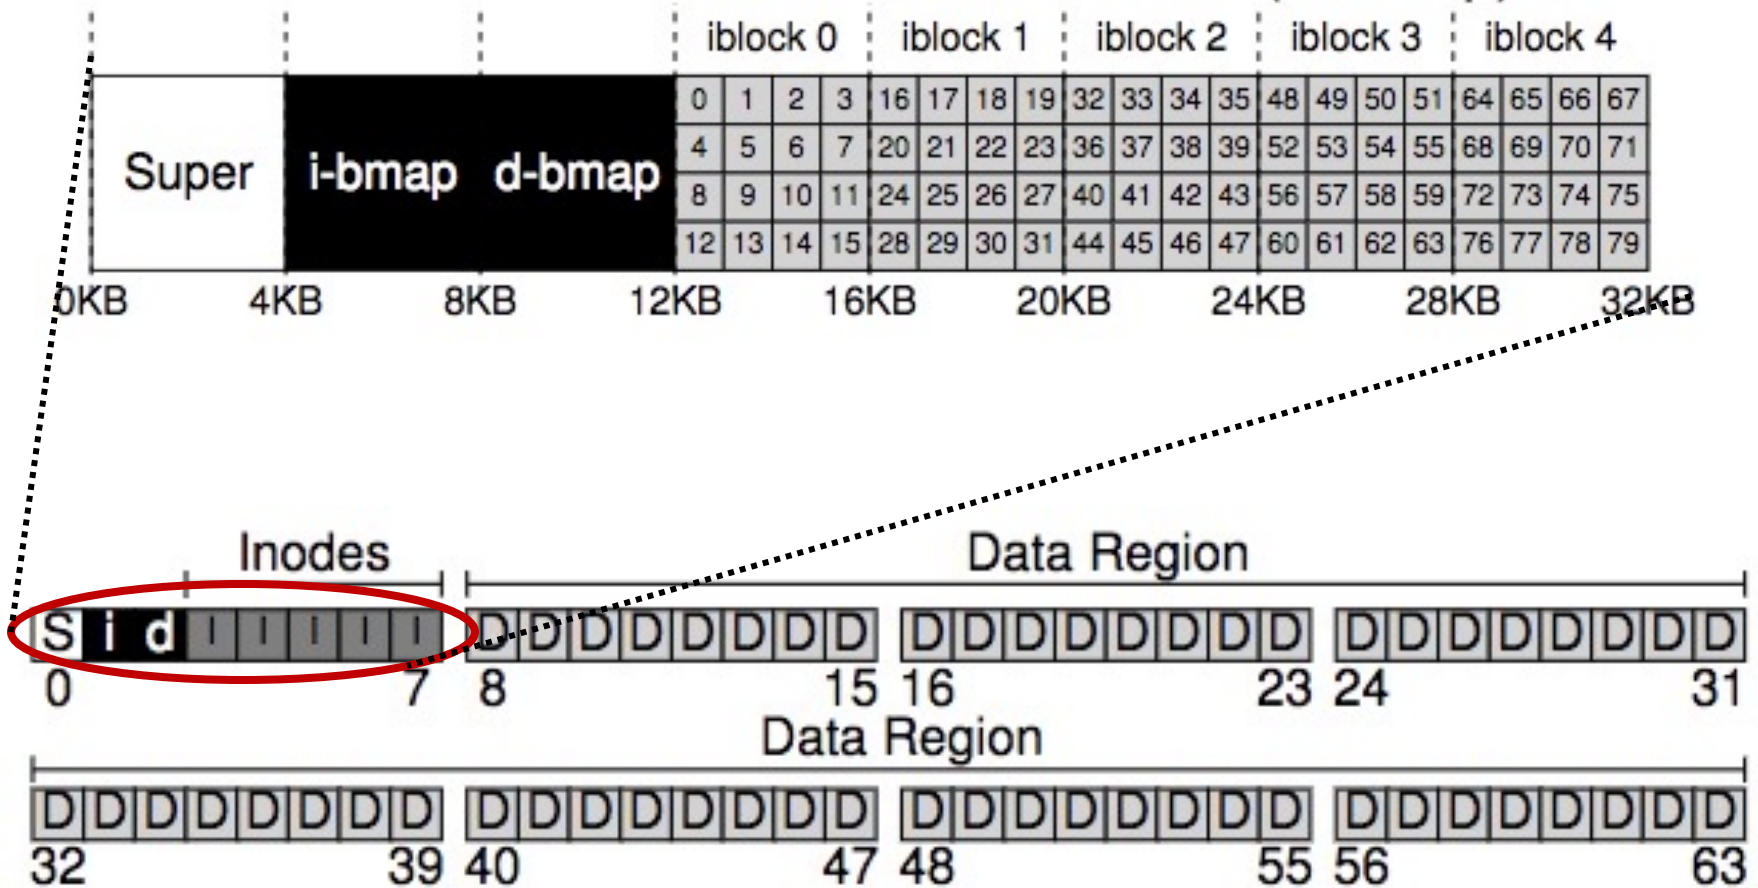


# On-Disk Structure: Superblock



# On-Disk Structure Overview

The Inode Table (Closeup)





# File System Operations

# Basic File System Operations

create /foo/bar

| data<br>bitmap | inode<br>bitmap | root<br>inode | foo<br>inode | bar<br>inode | root<br>data | foo<br>data |
|----------------|-----------------|---------------|--------------|--------------|--------------|-------------|
|                |                 |               |              |              |              |             |

# Basic File System Operations

create /foo/bar

[traverse]

| data<br>bitmap | inode<br>bitmap | root<br>inode | foo<br>inode | bar<br>inode | root<br>data | foo<br>data |
|----------------|-----------------|---------------|--------------|--------------|--------------|-------------|
|                |                 | read          |              |              | read         |             |

# Basic File System Operations

create /foo/bar

[traverse]

| data<br>bitmap | inode<br>bitmap | root<br>inode | foo<br>inode | bar<br>inode | root<br>data | foo<br>data |
|----------------|-----------------|---------------|--------------|--------------|--------------|-------------|
|                |                 | read          |              |              |              |             |
|                |                 |               | read         |              | read         |             |
|                |                 |               |              |              |              | read        |

# Basic File System Operations

create /foo/bar

[traverse]

| data<br>bitmap | inode<br>bitmap | root<br>inode | foo<br>inode | bar<br>inode | root<br>data | foo<br>data |
|----------------|-----------------|---------------|--------------|--------------|--------------|-------------|
|                |                 | read          |              |              | read         |             |
|                |                 |               | read         |              |              |             |
|                |                 |               |              |              |              | read        |

foo inode: we have permission  
foo data: bar doesn't exist

# Basic File System Operations

create /foo/bar

| data<br>bitmap | inode<br>bitmap | root<br>inode | foo<br>inode | bar<br>inode | root<br>data | foo<br>data |
|----------------|-----------------|---------------|--------------|--------------|--------------|-------------|
|                |                 | read          |              |              |              |             |
|                |                 |               | read         |              | read         |             |
|                |                 |               |              |              |              | read        |

# Basic File System Operations

create /foo/bar

[allocate inode]

| data<br>bitmap | inode<br>bitmap | root<br>inode | foo<br>inode | bar<br>inode | root<br>data | foo<br>data |
|----------------|-----------------|---------------|--------------|--------------|--------------|-------------|
|                |                 | read          |              |              |              |             |
|                |                 |               | read         |              | read         |             |
|                | read<br>write   |               |              |              |              | read        |

# Basic File System Operations

create /foo/bar

[populate inode]

| data<br>bitmap | inode<br>bitmap | root<br>inode | foo<br>inode | bar<br>inode  | root<br>data | foo<br>data |
|----------------|-----------------|---------------|--------------|---------------|--------------|-------------|
|                |                 | read          |              |               |              |             |
|                |                 |               | read         |               | read         |             |
|                | read<br>write   |               |              |               |              | read        |
|                |                 |               |              | read<br>write |              |             |



# Basic File System Operations

create /foo/bar

[add bar to /foo]

| data<br>bitmap | inode<br>bitmap | root<br>inode | foo<br>inode | bar<br>inode  | root<br>data | foo<br>data |
|----------------|-----------------|---------------|--------------|---------------|--------------|-------------|
|                |                 | read          |              |               |              |             |
|                |                 |               | read         |               | read         |             |
|                | read<br>write   |               |              |               |              | read        |
|                |                 |               |              | read<br>write |              |             |
|                |                 |               | write        |               |              |             |
|                |                 |               |              |               |              | write       |

# Basic File System Operations

write to /foo/bar

| data<br>bitmap | inode<br>bitmap | root<br>inode | foo<br>inode | bar<br>inode | root<br>data | foo<br>data | bar<br>data |
|----------------|-----------------|---------------|--------------|--------------|--------------|-------------|-------------|
|                |                 |               |              |              |              |             |             |



# Basic File System Operations

write to /foo/bar

[allocate block]

| data<br>bitmap | inode<br>bitmap | root<br>inode | foo<br>inode | bar<br>inode | root<br>data | foo<br>data | bar<br>data |
|----------------|-----------------|---------------|--------------|--------------|--------------|-------------|-------------|
| read<br>write  |                 | read          |              |              |              |             |             |

# Basic File System Operations

write to /foo/bar

[point to block]

| data<br>bitmap | inode<br>bitmap | root<br>inode | foo<br>inode | bar<br>inode | root<br>data | foo<br>data | bar<br>data |
|----------------|-----------------|---------------|--------------|--------------|--------------|-------------|-------------|
| read<br>write  |                 |               |              | read         |              |             |             |
|                |                 |               |              | write        |              |             |             |

# Basic File System Operations

write to /foo/bar

[point to block]

| data<br>bitmap | inode<br>bitmap | root<br>inode | foo<br>inode | bar<br>inode | root<br>data | foo<br>data | bar<br>data |
|----------------|-----------------|---------------|--------------|--------------|--------------|-------------|-------------|
| read<br>write  |                 |               |              | read         |              |             |             |
|                |                 |               |              | write        |              |             |             |
|                |                 |               |              |              |              |             | write       |

# Basic File System Operations

write to /foo/bar

[point to block]

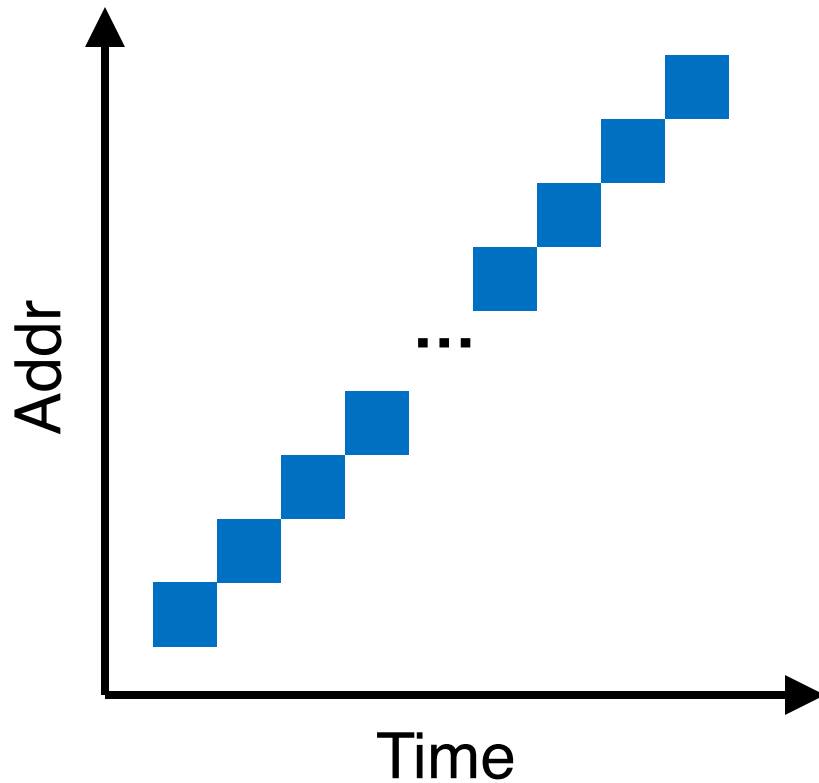
| data<br>bitmap | inode<br>bitmap | root<br>inode | foo<br>inode | bar<br>inode | root<br>data | foo<br>data | bar<br>data |
|----------------|-----------------|---------------|--------------|--------------|--------------|-------------|-------------|
|                |                 |               |              |              | dir blocks   |             | file        |
| read<br>write  |                 | read          |              |              |              |             |             |
|                |                 | write         |              |              |              |             |             |
|                |                 |               |              |              |              |             | write       |

# Locality & Data Layout

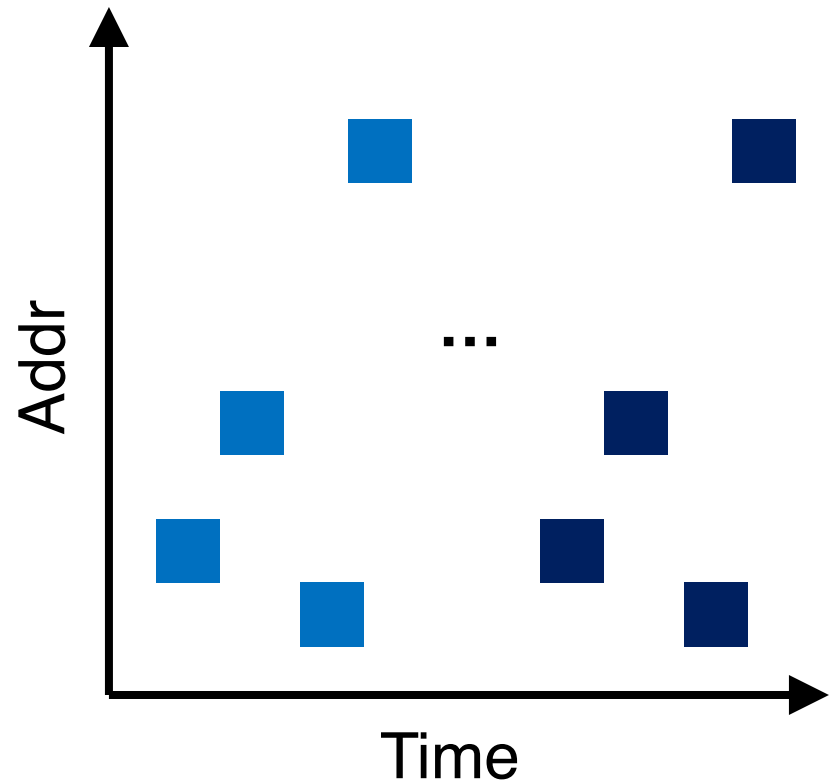


# Locality Types

Workload A

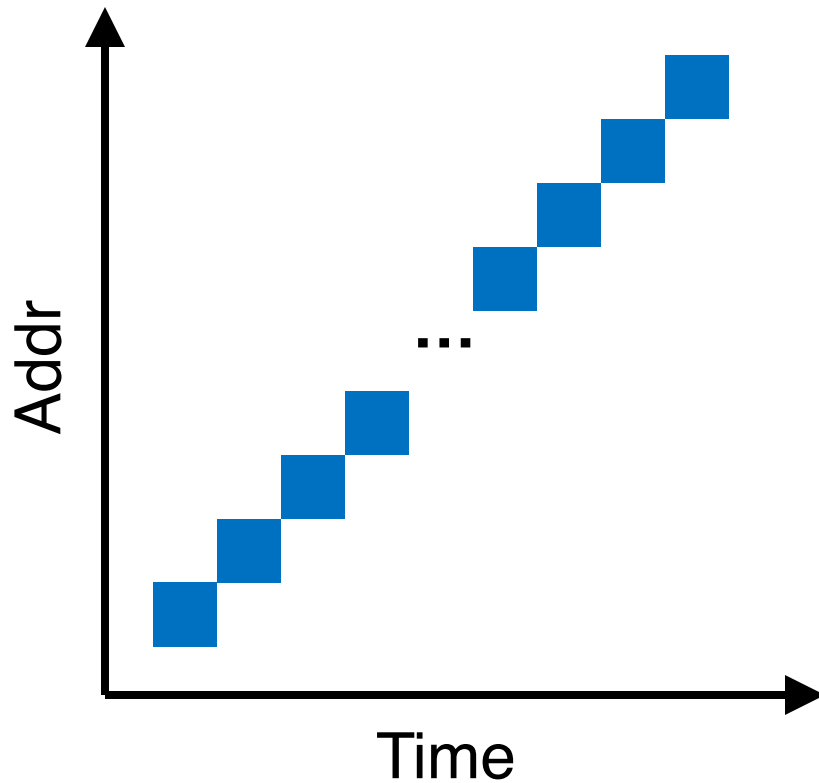


Workload B



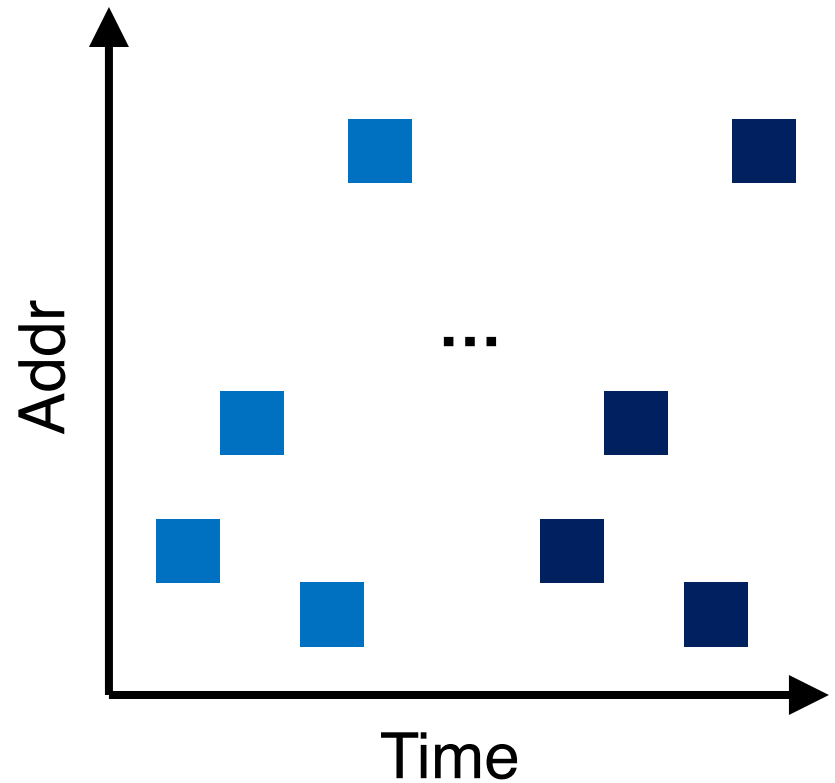
# Locality Types

Workload A



**Spatial Locality**

Workload B



**Temporal Locality**

# Locality Usefulness in the Context of Disk-based File Systems

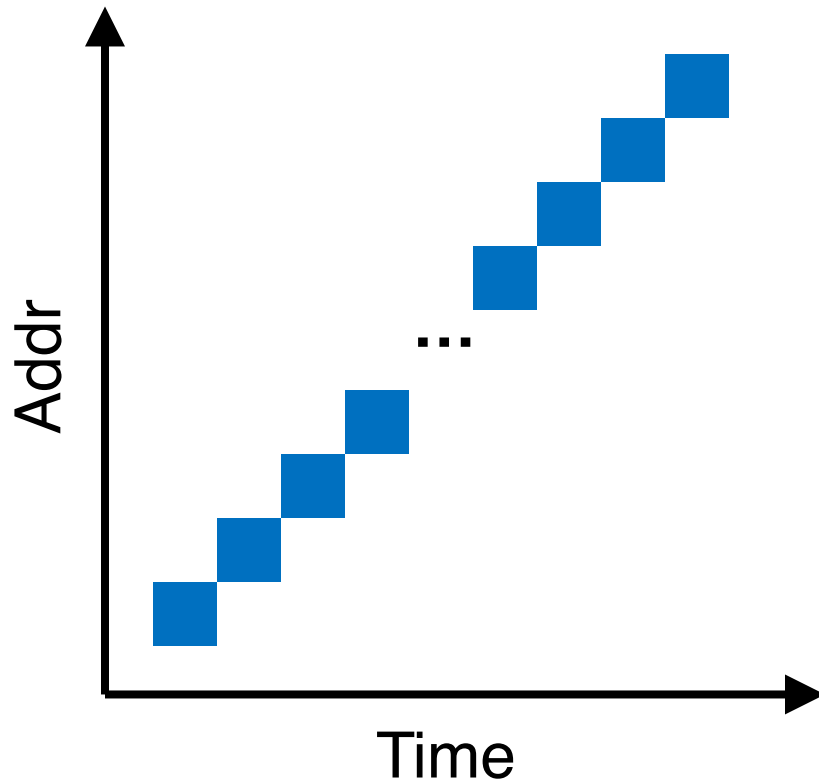
- What types of locality are useful for a [cache](#)?
- What types of locality are useful for a disk?

# Locality Usefulness in the Context of Disk-based File Systems

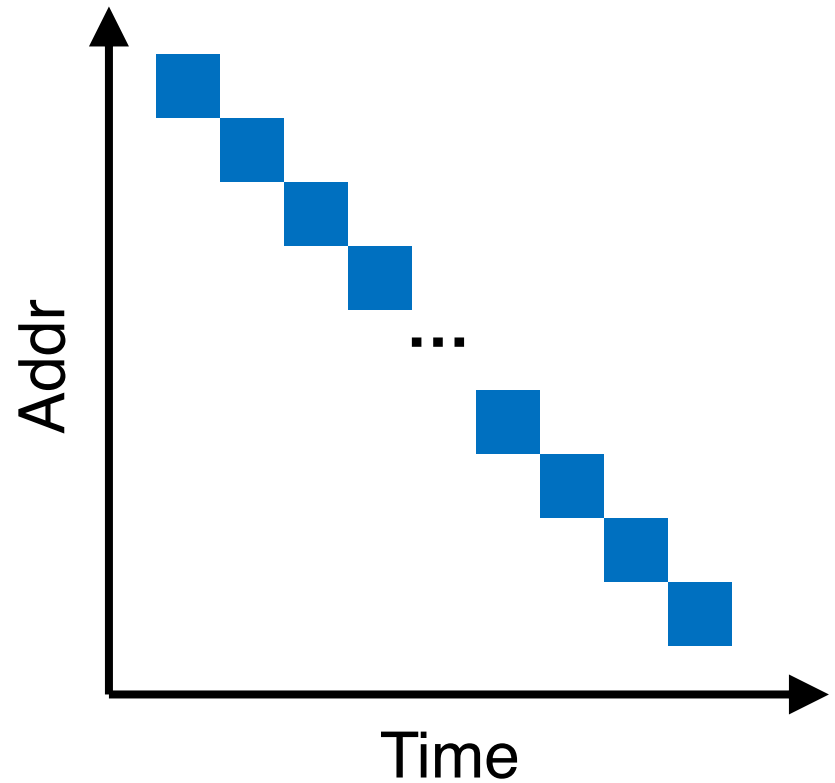
- What types of locality are useful for a **cache**?
  - Possibly, both spatial & temporal locality
- What types of locality are useful for a **disk**?
  - Spatial locality, since a disk sucks in random I/Os but can provide reasonably good sequential performance

# Order Matters Now for FS on Disk

Workload A

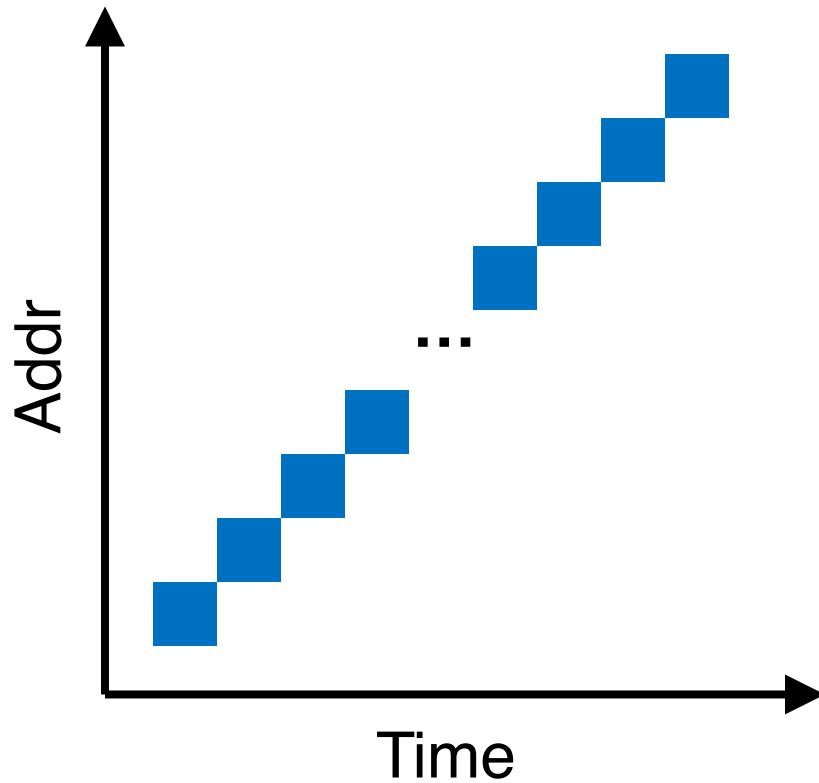


Workload B



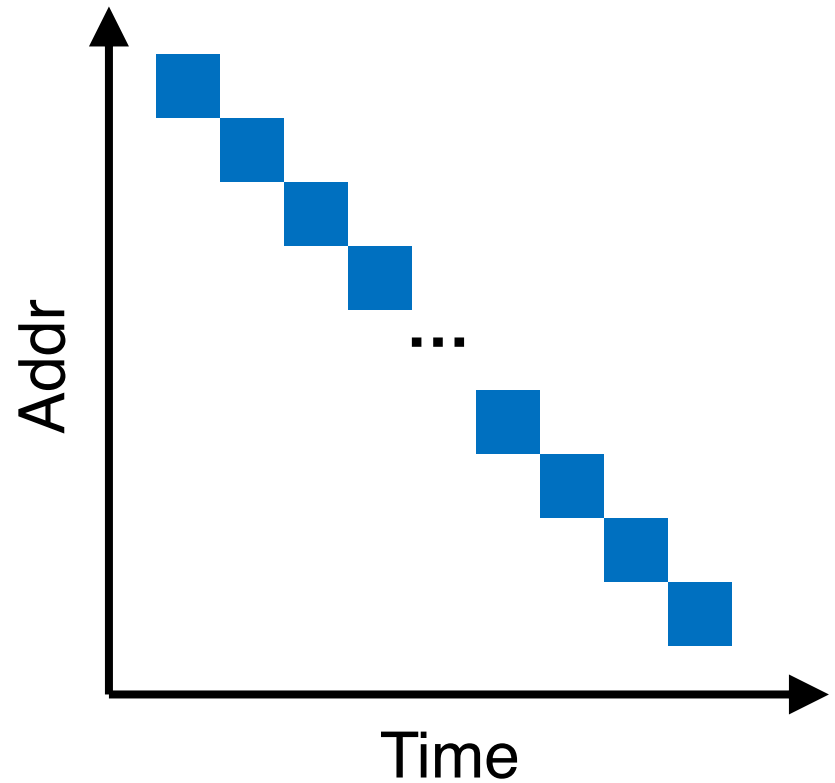
# Order Matters Now for FS on Disk

Workload A



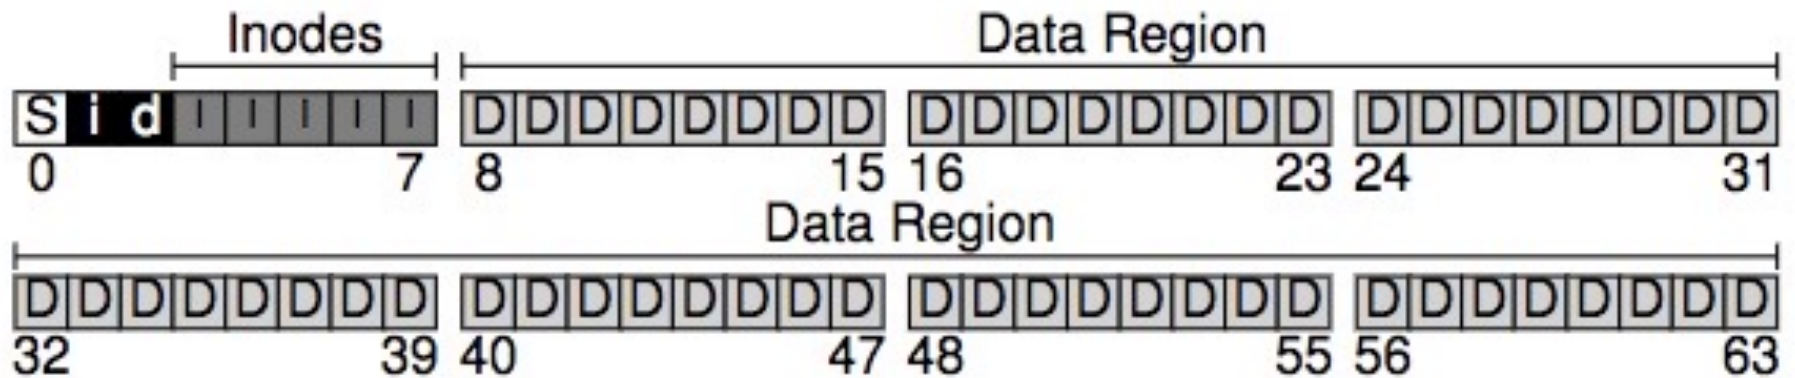
**Fast**

Workload B

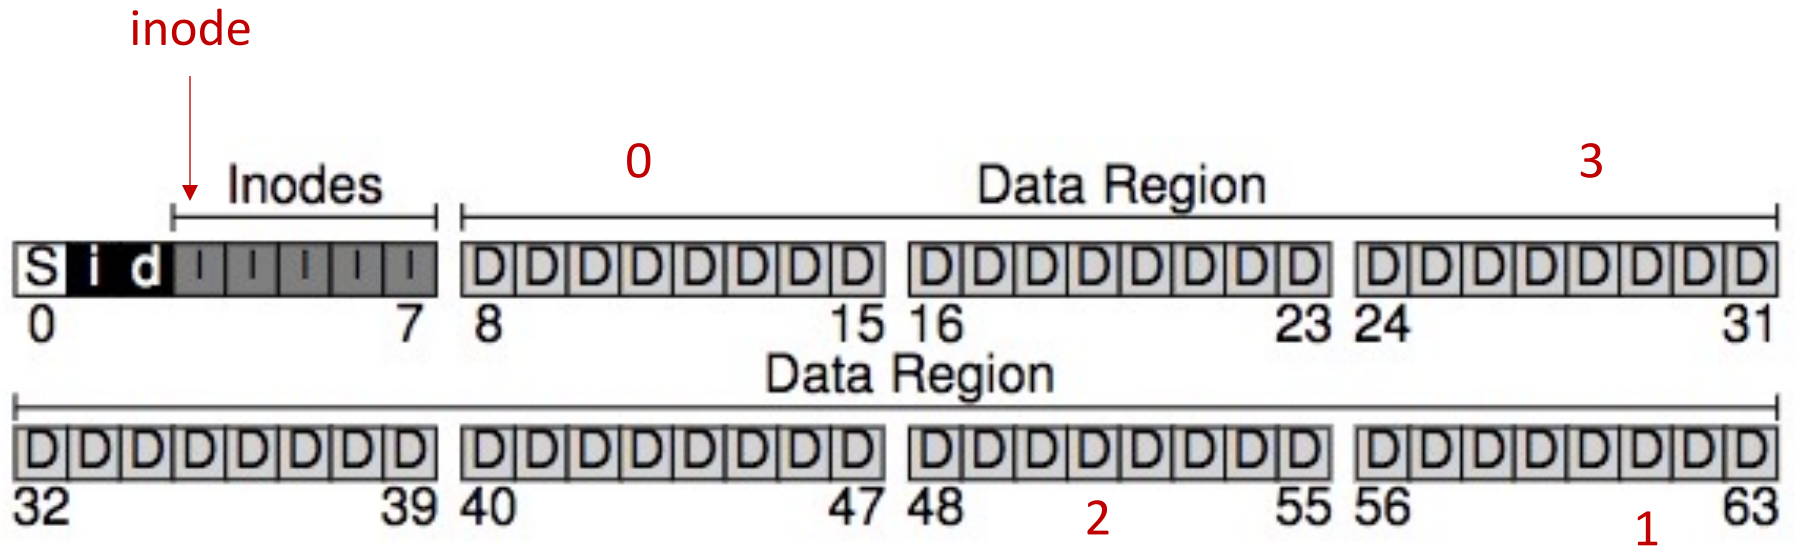


**Slow**

# Policy: Choose Inode, Data Blocks

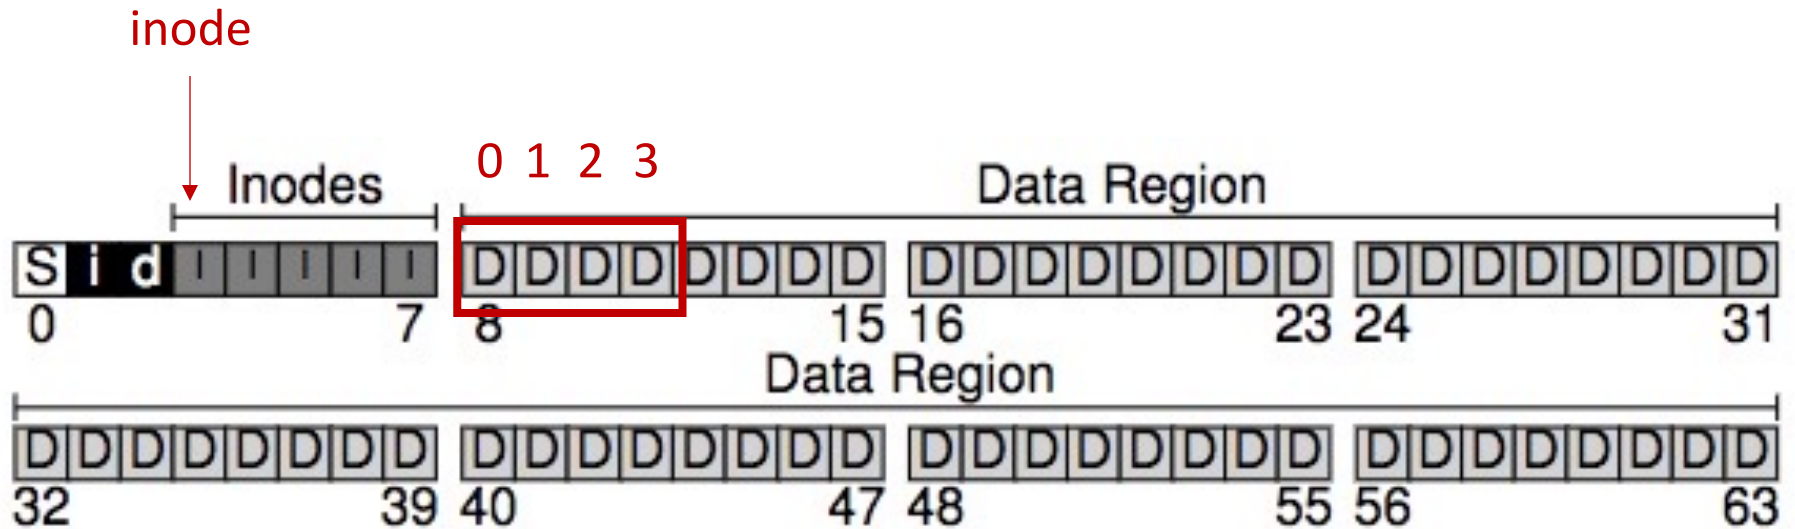


# Bad File Layout

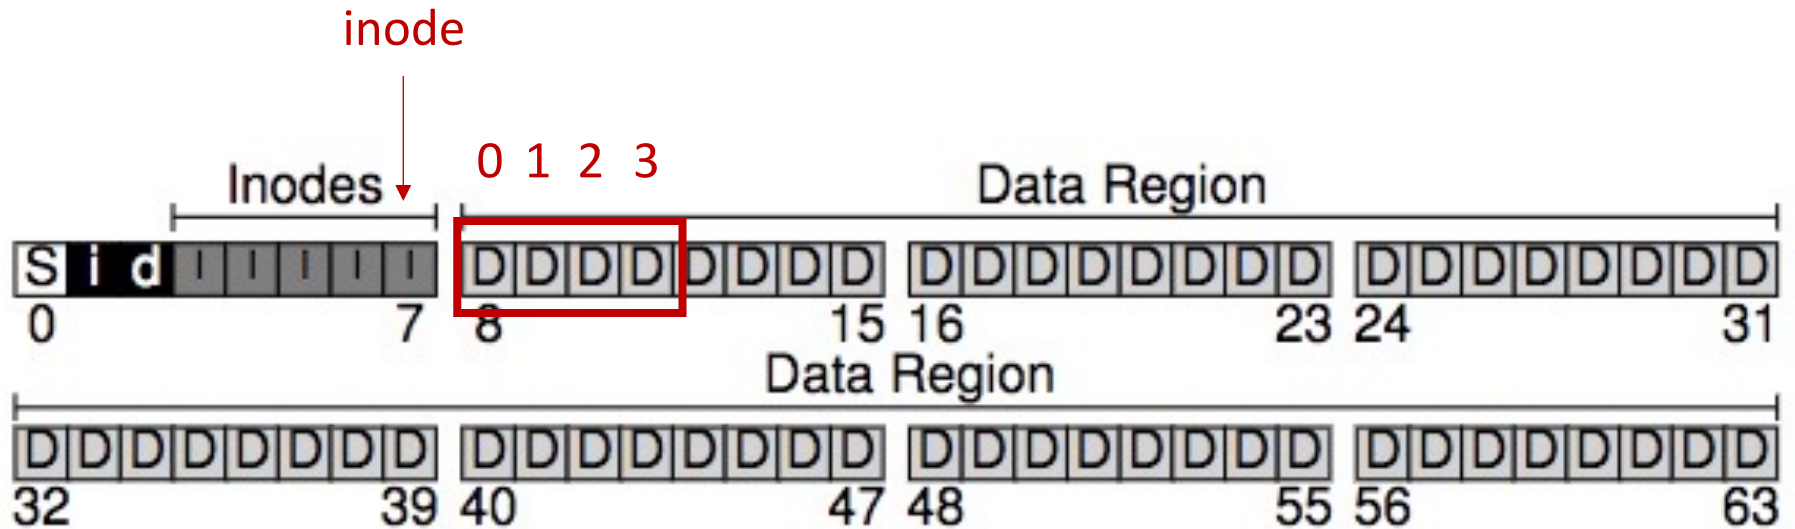




# Better File Layout



# Best File Layout



# Recap on Disks

# Properties of A Single Disk

- A single disk is slow
  - Kind of Okay sequential I/O performance
  - Really bad for random I/O

# Properties of A Single Disk

- A single disk is slow
  - Kind of Okay sequential I/O performance
  - Really bad for random I/O
- The storage capacity of a single disk is limited

# Properties of A Single Disk

- A single disk is slow
  - Kind of Okay sequential I/O performance
  - Really bad for random I/O
- The storage capacity of a single disk is limited
- A single disk is not reliable

# **RAID: Redundant Array of Inexpensive Disks**

# Wish List for a Disk

- Wish it to be faster
  - I/O is always the performance bottleneck



# Wish List for a Disk

- Wish it to be **faster**
  - I/O is always the performance bottleneck
- Wish it to be **larger**
  - More and more data needs to be stored

# Wish List for a Disk

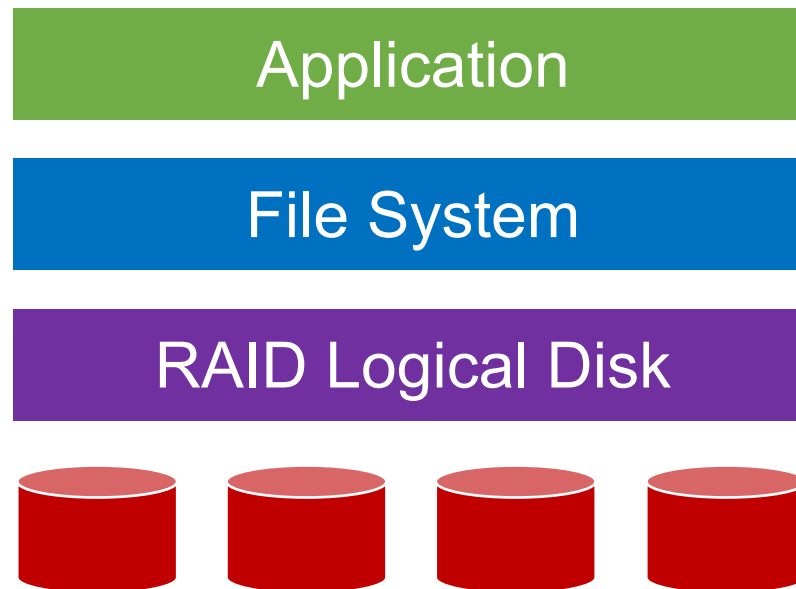
- Wish it to be **faster**
  - I/O is always the performance bottleneck
- Wish it to be **larger**
  - More and more data needs to be stored
- Wish it to be **more reliable**
  - We don't want our valuable data to be gone

# Only One Disk?

- Sometimes we want many disks
  - For higher performance
  - For larger capacity
  - For better reliability
- **Challenge:** Most file systems work on only one disk

# Solution: RAID

**RAID**: Redundant Array of Inexpensive Disks

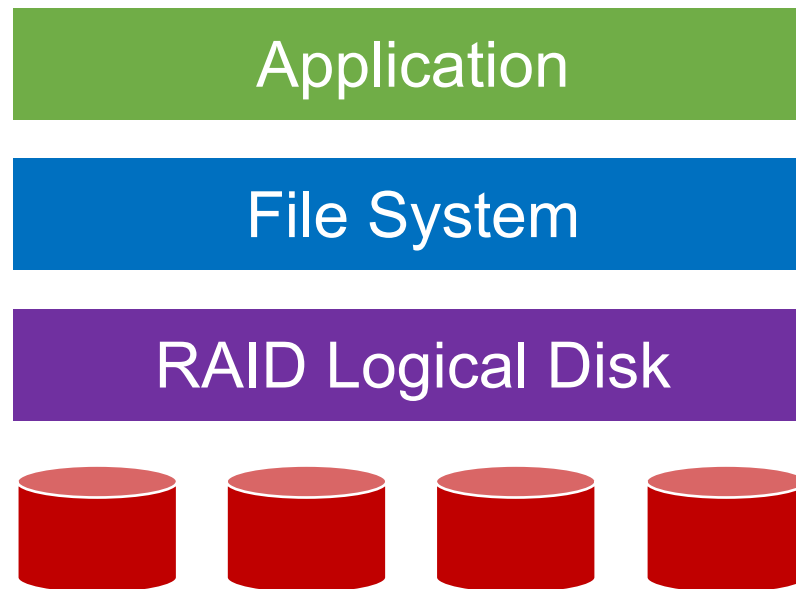


Build a logical disk from many physical disks

# Solution: RAID

**RAID**: Redundant Array of Inexpensive Disks

- RAID is
- Transparent
  - Deployable

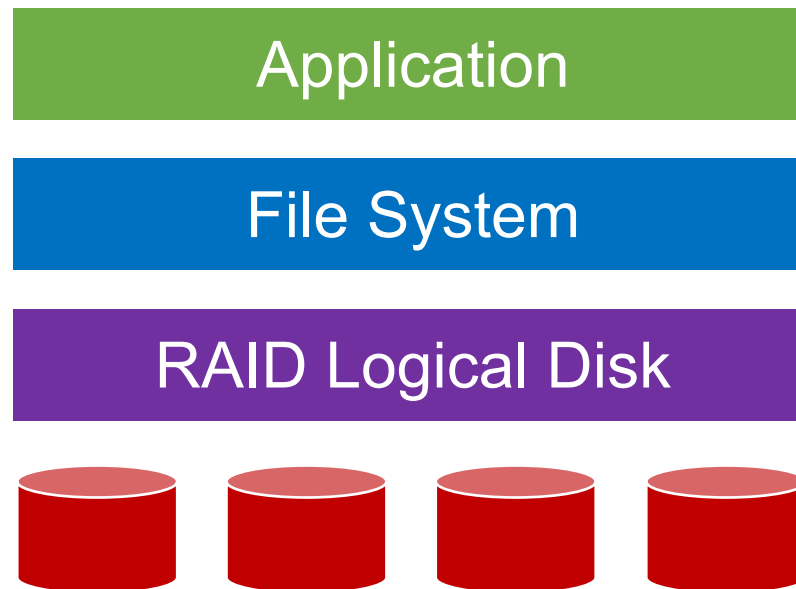


Build a logical disk from many physical disks

# Solution: RAID

**RAID**: Redundant Array of Inexpensive Disks

- RAID is
- Transparent
  - Deployable



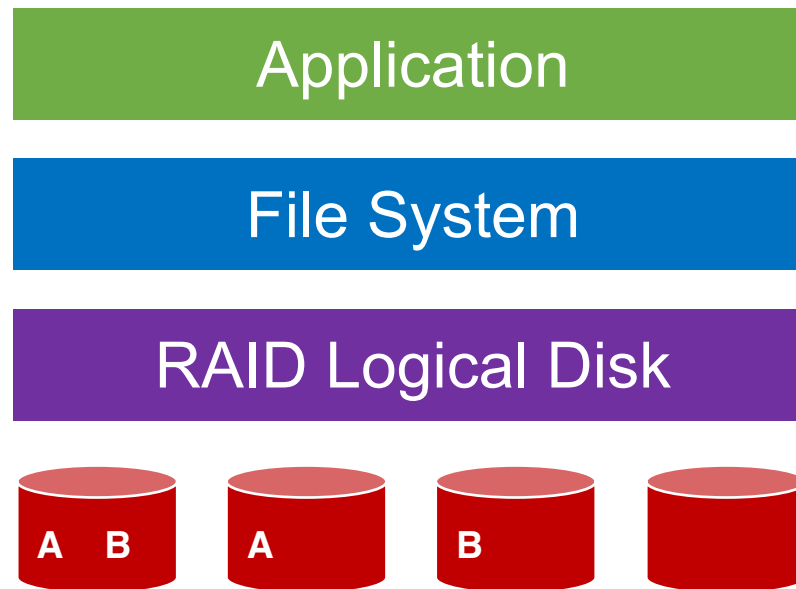
- Logical disks gives
- Performance
  - Capacity
  - Reliability

Build a logical disk from many physical disks

# Solution: RAID

**RAID**: Redundant Array of Inexpensive Disks

- RAID is
- Transparent
  - Deployable



- Logical disks gives
- Performance
  - Capacity
  - Reliability

Build a logical disk from many physical disks

# Why Inexpensive Disks?

- Economies of scale! Cheap disks are popular
- You can often get **many commodity** hardware components for the same price as a **few expensive** components

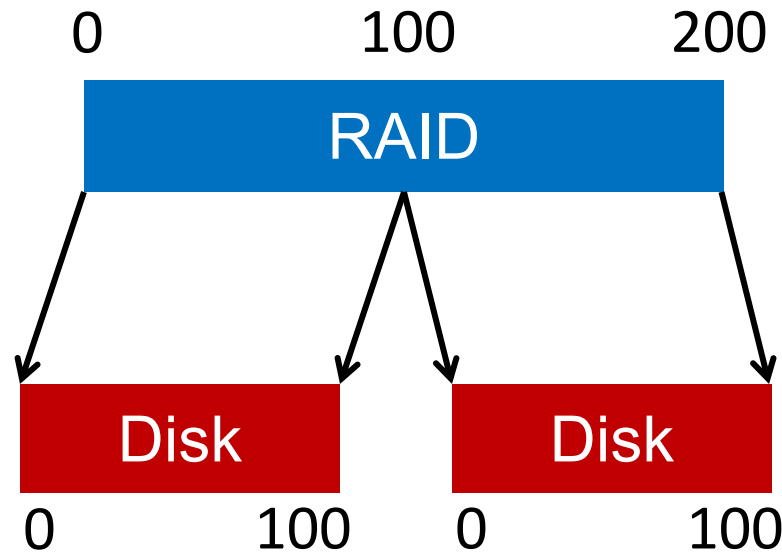


# Why Inexpensive Disks?

- Economies of scale! Cheap disks are popular
- You can often get **many commodity** hardware components for the same price as a **few expensive** components
- Strategy: Write software to **build high-quality logical devices from many cheap devices**
  - Tradeoff: To compensate poor properties of cheap devices

# General Strategy

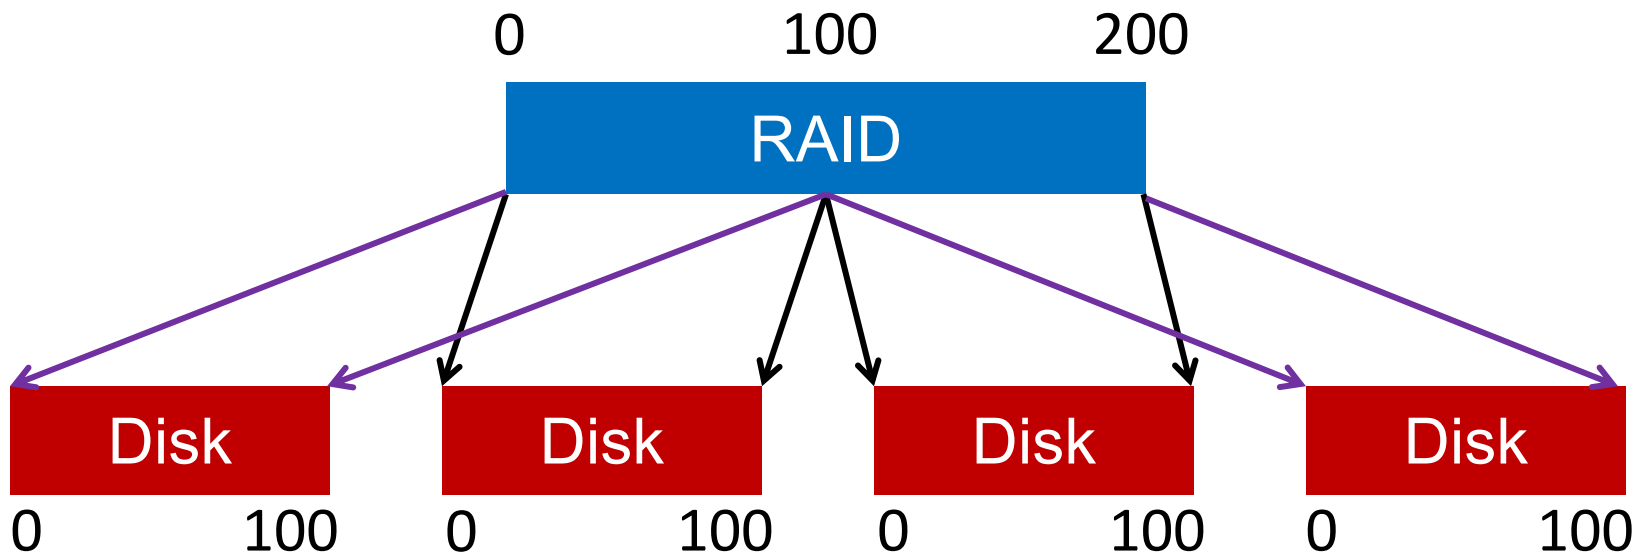
Build fast and large disks from smaller ones



# General Strategy

Build fast and large disks from smaller ones

Add more disks for **reliability++**!




# RAID Metrics

- Performance
  - How long does each workload take?
- Capacity
  - How much space can apps use?
- Reliability
  - How many disks can we safely lose?

# RAID Metrics

- Performance
  - How long does each workload take?
- Capacity
  - How much space can apps use?
- Reliability
  - How many disks can we safely lose?
  - Assume **fail-stop** model!


# RAID Levels



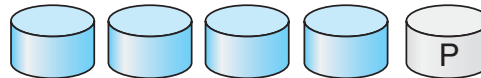
(a) RAID 0: non-redundant striping.




(b) RAID 1: mirrored disks.



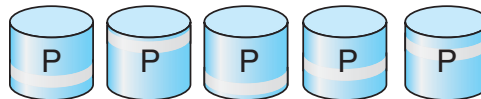
(c) RAID 2: memory-style error-correcting codes.



(d) RAID 3: bit-interleaved parity.




(e) RAID 4: block-interleaved parity.

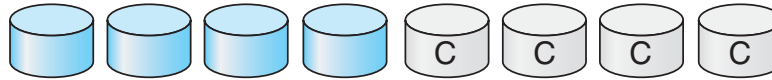


(f) RAID 5: block-interleaved distributed parity.


# RAID Level 0




(a) RAID 0: non-redundant striping.



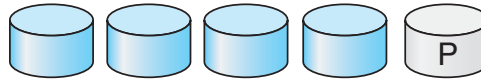
(b) RAID 1: mirrored disks.



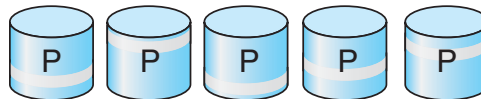
(c) RAID 2: memory-style error-correcting codes.



(d) RAID 3: bit-interleaved parity.



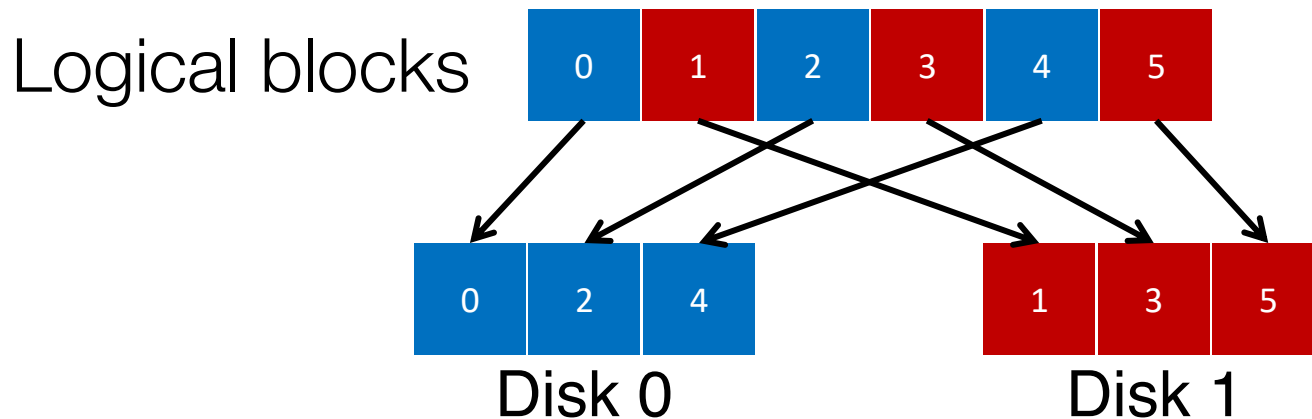
(e) RAID 4: block-interleaved parity.



(f) RAID 5: block-interleaved distributed parity.

# RAID-0: Striping

- No redundancy
- Serves as **upper bound** for
  - Performance
  - Capacity





# 4 Disks

| Disk 0 | Disk 1 | Disk 2 | Disk 3 |
|--------|--------|--------|--------|
| 0      | 1      | 2      | 3      |
| 4      | 5      | 6      | 7      |
| 8      | 9      | 10     | 11     |
| 12     | 13     | 14     | 15     |

# 4 Disks

|                | Disk 0 | Disk 1 | Disk 2 | Disk 3 |
|----------------|--------|--------|--------|--------|
|                | 0      | 1      | 2      | 3      |
| <b>stripe:</b> | 4      | 5      | 6      | 7      |
|                | 8      | 9      | 10     | 11     |
|                | 12     | 13     | 14     | 15     |

# How to Map?

- Given logical address A:
  - Disk = ...
  - Offset = ...

| Disk 0 | Disk 1 | Disk 2 | Disk 3 |
|--------|--------|--------|--------|
| 0      | 1      | 2      | 3      |
| 4      | 5      | 6      | 7      |
| 8      | 9      | 10     | 11     |
| 12     | 13     | 14     | 15     |

# How to Map?

- Given logical address A:
  - $\text{Disk} = A \% \text{disk\_count}$
  - $\text{Offset} = A / \text{disk\_count}$

| Disk 0 | Disk 1 | Disk 2 | Disk 3 |
|--------|--------|--------|--------|
| 0      | 1      | 2      | 3      |
| 4      | 5      | 6      | 7      |
| 8      | 9      | 10     | 11     |
| 12     | 13     | 14     | 15     |

# Mapping Example: Find Block 13

- Given logical address 13:
  - **Disk** =  $13 \% 4 = 1$
  - **Offset** =  $13 / 4 = 3$

|          | Disk 0 | Disk 1 | Disk 2 | Disk 3 |
|----------|--------|--------|--------|--------|
| Offset 0 | 0      | 1      | 2      | 3      |
| 1        | 4      | 5      | 6      | 7      |
| 2        | 8      | 9      | 10     | 11     |
| 3        | 12     | 13     | 14     | 15     |

# Chunk Size = 1

| Disk 0 | Disk 1 | Disk 2 | Disk 3 |
|--------|--------|--------|--------|
| 0      | 1      | 2      | 3      |
| 4      | 5      | 6      | 7      |
| 8      | 9      | 10     | 11     |
| 12     | 13     | 14     | 15     |

# Chunk Size = 1

| Disk 0 | Disk 1 | Disk 2 | Disk 3 |
|--------|--------|--------|--------|
| 0      | 1      | 2      | 3      |
| 4      | 5      | 6      | 7      |
| 8      | 9      | 10     | 11     |
| 12     | 13     | 14     | 15     |

# Chunk Size = 2

| Disk 0 | Disk 1 | Disk 2 | Disk 3 |                         |
|--------|--------|--------|--------|-------------------------|
| 0      | 2      | 4      | 6      | chunk size:<br>2 blocks |
| 1      | 3      | 5      | 7      |                         |
| 8      | 10     | 12     | 14     |                         |
| 9      | 11     | 13     | 15     |                         |

# Chunk Size = 1

| Disk 0 | Disk 1 | Disk 2 | Disk 3 |
|--------|--------|--------|--------|
| 0      | 1      | 2      | 3      |
| 4      | 5      | 6      | 7      |
| 8      | 9      | 10     | 11     |
| 12     | 13     | 14     | 15     |

In all following examples, we assume chunk size of 1

# Chunk Size = 2

| Disk 0 | Disk 1 | Disk 2 | Disk 3 |                         |
|--------|--------|--------|--------|-------------------------|
| 0      | 2      | 4      | 6      | chunk size:<br>2 blocks |
| 1      | 3      | 5      | 7      |                         |
| 8      | 10     | 12     | 14     |                         |
| 9      | 11     | 13     | 15     |                         |



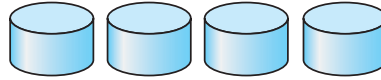
# RAID-0 Analysis

1. What is capacity?
2. How many disks can fail?
3. Throughput?
4. Latency?

# RAID-0 Analysis

1. What is capacity?  $N * C$
2. How many disks can fail? 0
3. Throughput?  $N * S$  and  $N * R$
4. Latency?  $D$

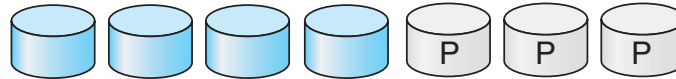
# RAID Level 1



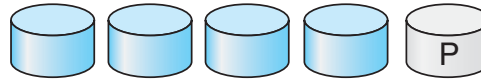
(a) RAID 0: non-redundant striping.



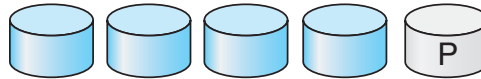
(b) RAID 1: mirrored disks.




(c) RAID 2: memory-style error-correcting codes.



(d) RAID 3: bit-interleaved parity.



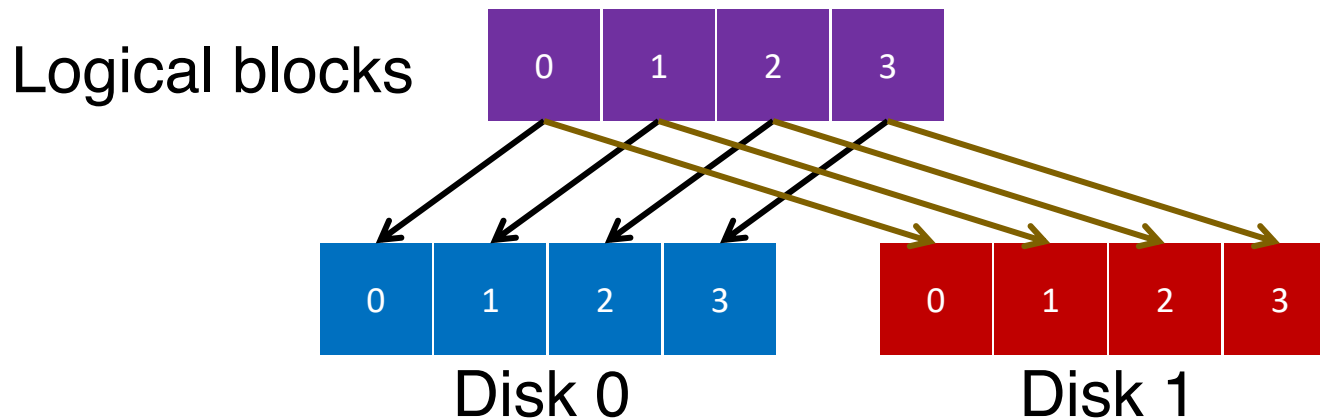
(e) RAID 4: block-interleaved parity.



(f) RAID 5: block-interleaved distributed parity.

# RAID-1: Mirroring

- RAID-1 keeps two copies of each block



# Assumption

- Assume disks are **fail-stop**
  - Two states
    - They work or they don't
  - We know when they don't work

# 4 Disks

| Disk 0 | Disk 1 | Disk 2 | Disk 3 |
|--------|--------|--------|--------|
| 0      | 0      | 1      | 1      |
| 2      | 2      | 3      | 3      |
| 4      | 4      | 5      | 5      |
| 6      | 6      | 7      | 7      |

# 4 Disks

| Disk 0 | Disk 1 | Disk 2 | Disk 3 |
|--------|--------|--------|--------|
| 0      | 0      | 1      | 1      |
| 2      | 2      | 3      | 3      |
| 4      | 4      | 5      | 5      |
| 6      | 6      | 7      | 7      |

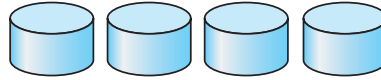
**How many disks can fail?**

# RAID-1 Analysis

1. What is capacity?  $N/2 * C$
2. How many disks can fail? 1 or maybe  $N / 2$
3. Throughput?
  - Seq read:  $N/2 * S$
  - Seq write:  $N/2 * S$
  - Rand read:  $N * R$
  - Rand write:  $N/2 * R$
4. Latency?  $D$



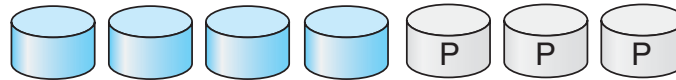
# RAID Level 4



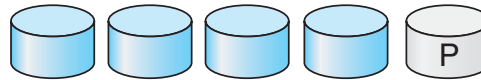
(a) RAID 0: non-redundant striping.



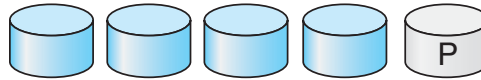
(b) RAID 1: mirrored disks.



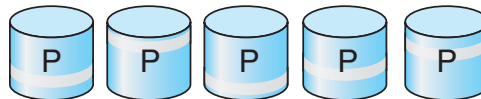
(c) RAID 2: memory-style error-correcting codes.



(d) RAID 3: bit-interleaved parity.




(e) RAID 4: block-interleaved parity.




(f) RAID 5: block-interleaved distributed parity.


# RAID-4



# RAID-4



# RAID-4



# RAID-4: Strategy

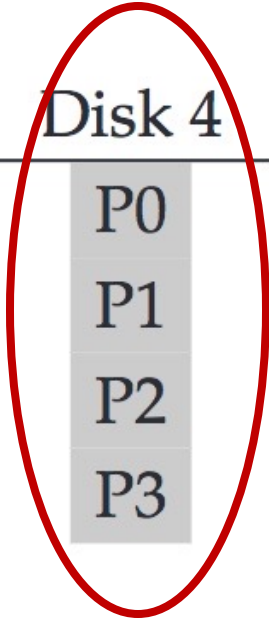
- Use **parity** disk
- In algebra, if an **equation** has  $N$  variables, and  $N-1$  are known, you can also solve for the unknown
- Treat the sectors/blocks across disks in a stripe as an equation

# RAID-4: Strategy

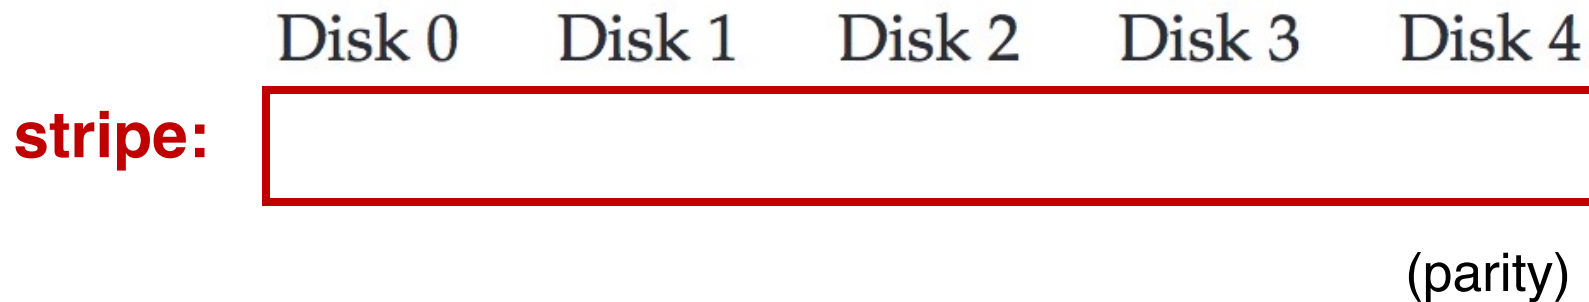
- Use **parity** disk
- In algebra, if an **equation** has  $N$  variables, and  $N-1$  are known, you can also solve for the unknown
- Treat the sectors/blocks across disks in a stripe as an equation
- A **failed disk** is like an unknown **in that equation**

# 5 Disks

| Disk 0 | Disk 1 | Disk 2 | Disk 3 | Disk 4 |
|--------|--------|--------|--------|--------|
| 0      | 1      | 2      | 3      | P0     |
| 4      | 5      | 6      | 7      | P1     |
| 8      | 9      | 10     | 11     | P2     |
| 12     | 13     | 14     | 15     | P3     |



# Example





# Example

|                | Disk 0 | Disk 1 | Disk 2 | Disk 3 | Disk 4 |
|----------------|--------|--------|--------|--------|--------|
| <b>stripe:</b> | 4      | 3      | 0      | 2      |        |

(parity)

# Example

|                | Disk 0 | Disk 1 | Disk 2 | Disk 3 | Disk 4 |
|----------------|--------|--------|--------|--------|--------|
| <b>stripe:</b> | 4      | 3      | 0      | 2      | 9      |

(parity)

# Example

|                | Disk 0 | Disk 1 | Disk 2 | Disk 3 | Disk 4 |
|----------------|--------|--------|--------|--------|--------|
| <b>stripe:</b> | X      | 3      | 0      | 2      | 9      |

(parity)

# Example

|                | Disk 0   | Disk 1 | Disk 2 | Disk 3 | Disk 4 |
|----------------|----------|--------|--------|--------|--------|
| <b>stripe:</b> | <b>4</b> | 3      | 0      | 2      | 9      |

(parity)

# Parity Function: XOR Example

| C0 | C1 | C2 | C3 | P                         |
|----|----|----|----|---------------------------|
| 0  | 0  | 1  | 1  | $\text{XOR}(0,0,1,1) = 0$ |
| 0  | 1  | 0  | 0  | $\text{XOR}(0,1,0,0) = 1$ |

# Parity Function: XOR Example

| C0 | C1 | C2 | C3 | P                         |
|----|----|----|----|---------------------------|
| 0  | 0  | 1  | 1  | $\text{XOR}(0,0,1,1) = 0$ |
| 0  | 1  | 0  | 0  | $\text{XOR}(0,1,0,0) = 1$ |

XOR function:

- $P = 0$ : The number of 1 in a stripe must be an even number
- $P = 1$ : The number of 1 in a stripe must be an odd number

# Parity Function: XOR Example

|                | Block0 | Block1 | Block2 | Block3 | Parity |
|----------------|--------|--------|--------|--------|--------|
| <b>stripe:</b> | 00     | 10     | 11     | 10     | 11     |
|                | 10     | 01     | 00     | 01     | 10     |

XOR function:

- $P = 0$ : The number of 1 in a stripe must be an even number
- $P = 1$ : The number of 1 in a stripe must be an odd number

# Parity Function: XOR Example

|                | Block0                               | Block1                      | Block2                      | Block3                      | Parity                      |
|----------------|--------------------------------------|-----------------------------|-----------------------------|-----------------------------|-----------------------------|
| <b>stripe:</b> | <div><div>X</div><div>10</div></div> | <div>10</div> <div>01</div> | <div>11</div> <div>00</div> | <div>10</div> <div>01</div> | <div>11</div> <div>10</div> |

XOR function:

- $P = 0$ : The number of 1 in a stripe must be an even number
- $P = 1$ : The number of 1 in a stripe must be an odd number



# Parity Function: XOR Example

|                | Block0                               | Block1                      | Block2                      | Block3                      | Parity                      |
|----------------|--------------------------------------|-----------------------------|-----------------------------|-----------------------------|-----------------------------|
| <b>stripe:</b> | <div><div>X</div><div>10</div></div> | <div>10</div> <div>01</div> | <div>11</div> <div>00</div> | <div>10</div> <div>01</div> | <div>11</div> <div>10</div> |

$$\text{Block0} = \text{XOR}(10, 11, 10, 11) = 00$$

XOR function:

- $P = 0$ : The number of 1 in a stripe must be an even number
- $P = 1$ : The number of 1 in a stripe must be an odd number

# Parity Function: XOR Example

|                | Block0 | Block1 | Block2 | Block3 | Parity |
|----------------|--------|--------|--------|--------|--------|
| <b>stripe:</b> | 00     | 10     | 11     | 10     | 11     |
|                | 10     | 01     | 00     | 01     | 10     |

$$\text{Block0} = \text{XOR}(10, 11, 10, 11) = 00$$

XOR function:

- $P = 0$ : The number of 1 in a stripe must be an even number
- $P = 1$ : The number of 1 in a stripe must be an odd number

# RAID-4 Analysis

1. What is capacity?  $(N-1) * C$
2. How many disks can fail? 1
3. Throughput?
  - Seq read:  $(N-1) * S$
  - Seq write:  $(N-1) * S$
  - Rand read:  $(N-1) * R$
  - Rand write:  $R/2$
4. Latency?  $D, 2D$

# RAID-4 Analysis: Random Write


Random write to 4, 13, and respective parity blocks

| Disk 0 | Disk 1 | Disk 2 | Disk 3 | Disk 4 |
|--------|--------|--------|--------|--------|
| 0      | 1      | 2      | 3      | P0     |
| *4     | 5      | 6      | 7      | +P1    |
| 8      | 9      | 10     | 11     | P2     |
| 12     | *13    | 14     | 15     | +P3    |


**Small write problem** (for parity-based RAID):

Parity disk serializes all random writes; each **logical** I/O generates two **physical** I/Os (**one read and one write for parity P1**)


# RAID Level 5




(a) RAID 0: non-redundant striping.




(b) RAID 1: mirrored disks.




(c) RAID 2: memory-style error-correcting codes.



(d) RAID 3: bit-interleaved parity.



(e) RAID 4: block-interleaved parity.



(f) RAID 5: block-interleaved distributed parity.

# RAID-5: Rotating Parity

| Disk 0 | Disk 1 | Disk 2 | Disk 3 | Disk 4 |
|--------|--------|--------|--------|--------|
| 0      | 1      | 2      | 3      | P0     |
| 5      | 6      | 7      | P1     | 4      |
| 10     | 11     | P2     | 8      | 9      |
| 15     | P3     | 12     | 13     | 14     |
| P4     | 16     | 17     | 18     | 19     |

RAID-5 works almost identically to RAID-4, except that it rotates the parity block across drives

# RAID-5 Analysis

1. What is capacity?  $(N-1) * C$
2. How many disks can fail? 1
3. Throughput?
  - Seq read:  $(N-1) * S$
  - Seq write:  $(N-1) * S$
  - Rand read:  $N * R$
  - Rand write: ???
4. Latency?  $D, 2D$

# RAID-5: Random Write

| Write  |        |        |        |        |
|--------|--------|--------|--------|--------|
| Disk 0 | Disk 1 | Disk 2 | Disk 3 | Disk 4 |
| 0      | 1      | 2      | 3      | P0     |
| 5      | 6      | 7      | P1     | 4      |
| 10     | 11     | P2     | 8      | 9      |
| 15     | P3     | 12     | 13     | 14     |
| P4     | 16     | 17     | 18     | 19     |

Random write to Block 10 on Disk 0



# RAID-5: Random Write

1. Read

| Disk 0 | Disk 1 | Disk 2 | Disk 3 | Disk 4 |
|--------|--------|--------|--------|--------|
| 0      | 1      | 2      | 3      | P0     |
| 5      | 6      | 7      | P1     | 4      |
| 10     | 11     | P2     | 8      | 9      |
| 15     | P3     | 12     | 13     | 14     |
| P4     | 16     | 17     | 18     | 19     |

Random write to Block 10 on Disk 0

1. Read Block 10

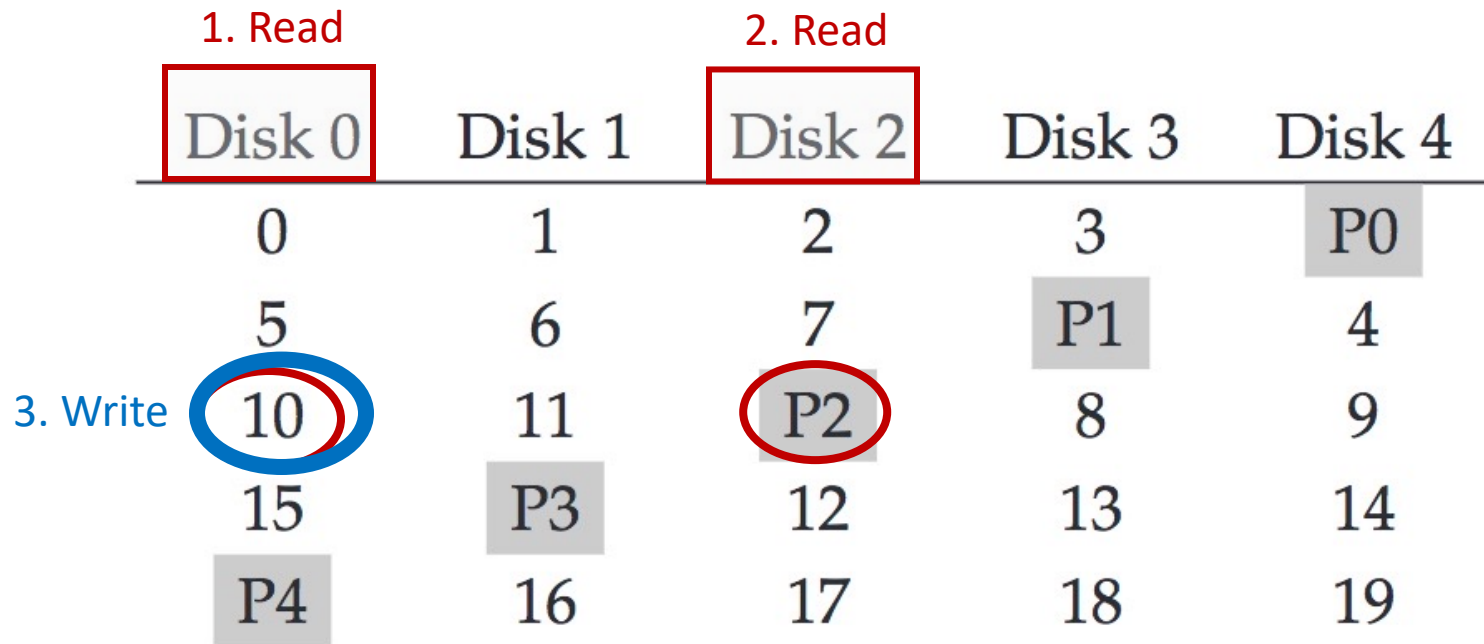
# RAID-5: Random Write

| 1. Read |    | 2. Read |    |    |
|---------|----|---------|----|----|
| Disk 0  |    | Disk 2  |    |    |
| 0       | 1  | 2       | 3  | P0 |
| 5       | 6  | 7       | P1 | 4  |
| 10      | 11 | P2      | 8  | 9  |
| 15      | P3 | 12      | 13 | 14 |
| P4      | 16 | 17      | 18 | 19 |

Random write to Block 10 on Disk 0

1. Read Block 10
2. Read the Parity P2

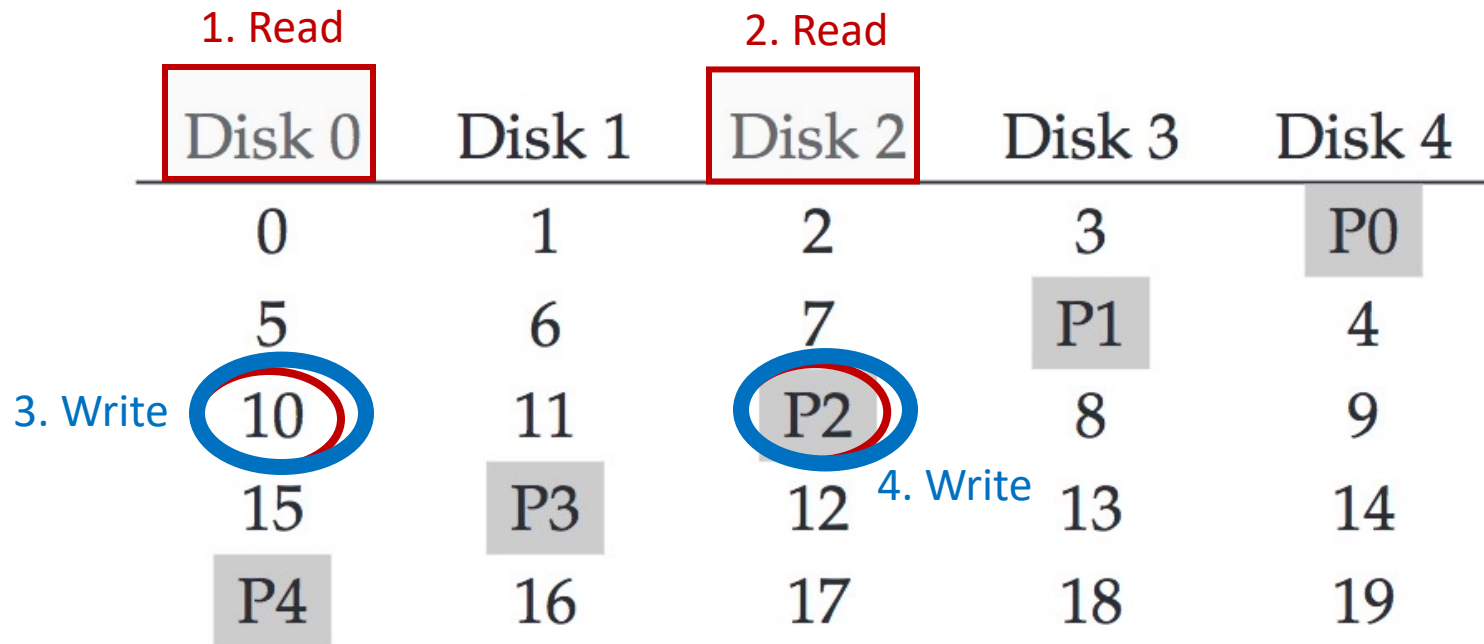
# RAID-5: Random Write



Random write to Block 10 on Disk 0

1. Read Block 10
2. Read the Parity P2
3. Write new data in Block 10

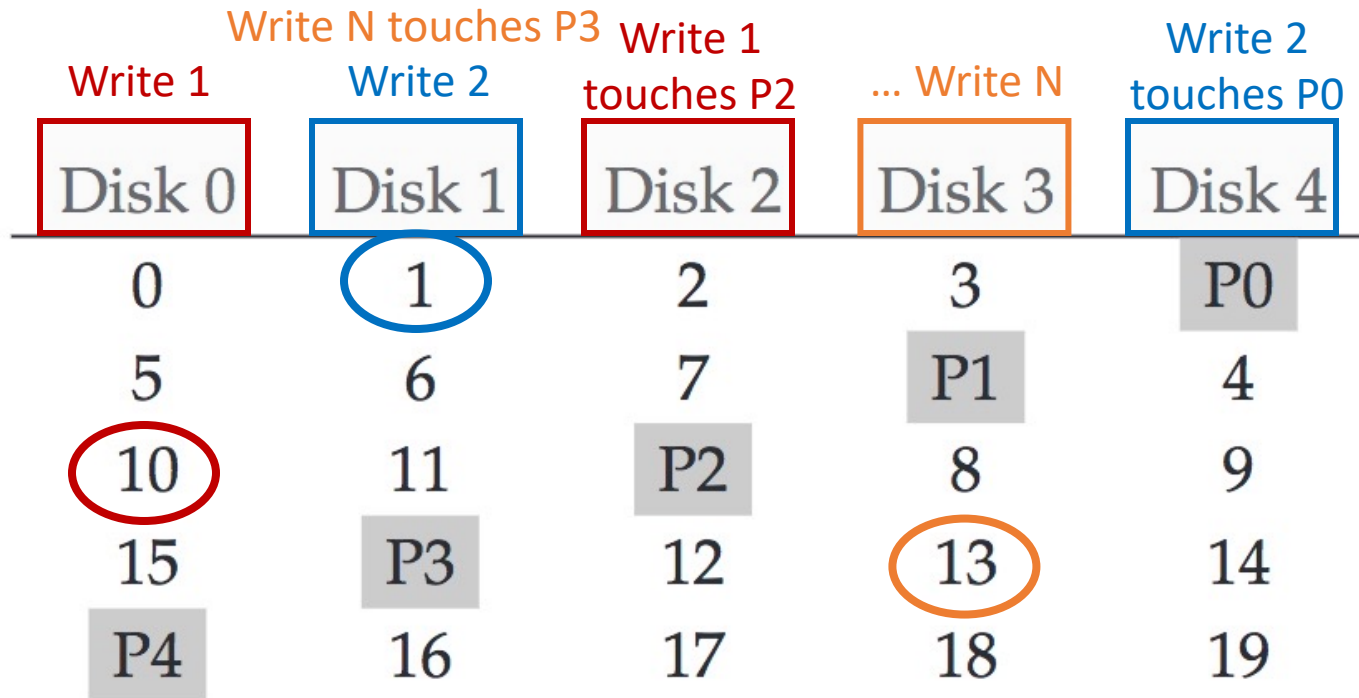
# RAID-5: Random Write



## Random write to Block 10 on Disk 0

1. Read Block 10
2. Read the Parity P2
3. Write new data in Block 10
4. Write new parity P2

# RAID-5: Random Write



Performance reasoning

Generally, for a large number of random read/write requests, RAID-5 will be able to keep all disks busy: thus  $N * R$



Each random (RAID-5) writes generates 4 physical I/O operations: thus  $N * R / 4$

# RAID-5 Analysis

1. What is capacity?  $(N-1) * C$
2. How many disks can fail? 1
3. Throughput?
  - Seq read:  $(N-1) * S$
  - Seq write:  $(N-1) * S$
  - Rand read:  $N * R$
  - Rand write:  $N * R/4$
4. Latency?  $D, 2D$

# Summary: All RAID's

|        | Reliability | Capacity    |
|--------|-------------|-------------|
| RAID-0 | 0           | $C * N$     |
| RAID-1 | 1 or $N/2$  | $C * N/2$   |
| RAID-4 | 1           | $C * (N-1)$ |
| RAID-5 | 1           | $C * (N-1)$ |

# Summary: All RAID's

|        | Seq Read    | Seq Write   | Rand Read   | Rand Write |
|--------|-------------|-------------|-------------|------------|
| RAID-0 | $N * S$     | $N * S$     | $N * R$     | $N * R$    |
| RAID-1 | $N/2 * S$   | $N/2 * S$   | $N * R$     | $N/2 * R$  |
| RAID-4 | $(N-1) * S$ | $(N-1) * S$ | $(N-1) * R$ | $R/2$      |
| RAID-5 | $(N-1) * S$ | $(N-1) * S$ | $N * R$     | $N/4 * R$  |



# Please Read the Textbook!

Do read the text chapter “RAID”: it has in-depth discussion of the various performance analyses covered in lecture.